

Forecasting monthly sales

January 9, 2020

1 Predict Future Sales

This notebook was completed as part of the Coursera Advanced Machine Learning specialisation as the final project for course 2.

This competition involved working with a challenging time-series dataset consisting of daily sales data, provided by one of the largest Russian software firms - 1C Company.

The task was to predict total sales for every product and store in the next month (November 2015).

The dataset can be accessed here: <https://www.kaggle.com/c/competitive-data-science-predict-future-sales/overview>

```
[2]: import numpy as np
import pandas as pd
import os
from itertools import product
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.ensemble import RandomForestRegressor
import xgboost as xgb
from xgboost import plot_importance
import lightgbm as lgb
import seaborn as sns
import matplotlib.pyplot as plt
import time
import sys
import gc
import pickle
from tqdm import tqdm
import warnings
warnings.filterwarnings("ignore")

%matplotlib inline

Validation = False
reduce_size = False
#num_first_level_models = 3
```

```

SEED = 0
start_time = time.time()

pd.set_option('display.max_rows', 99)
pd.set_option('display.max_columns', 50)

def downcast_dtypes(df):
    float_cols = [c for c in df if df[c].dtype == "float64"]
    int_cols = [c for c in df if df[c].dtype in ["int64", "int32"]]
    df[float_cols] = df[float_cols].astype(np.float32)
    df[int_cols] = df[int_cols].astype(np.int16)
    return df

#function to calculate RMSE
def rmse(actual, predictions):
    return np.sqrt(mean_squared_error(actual, predictions))

```

/Users/charlottefettes/opt/anaconda3/lib/python3.7/site-packages/lightgbm/__init__.py:48: UserWarning: Starting from version 2.2.1, the library file in distribution wheels for macOS is built by the Apple Clang (Xcode_8.3.3) compiler.

This means that in case of installing LightGBM from PyPI via the ``pip install lightgbm`` command, you don't need to install the gcc compiler anymore. Instead of that, you need to install the OpenMP library, which is required for running LightGBM on the system with the Apple Clang compiler. You can install the OpenMP library by the following command: ``brew install libomp``.

"You can install the OpenMP library by the following command: ``brew install libomp``.", UserWarning)

```

[258]: #read data into notebook
items = pd.read_csv('items.csv')
shops = pd.read_csv('shops.csv')
cats = pd.read_csv('item_categories.csv')
train = pd.read_csv('sales_train.csv')
test = pd.read_csv('test.csv')

```

1.0.1 Test set analysis

```

[3]: good_sales = test.merge(train, on=['item_id', 'shop_id'], how='left').dropna()
good_pairs = test[test['ID'].isin(good_sales['ID'])]
others = test[~(test['ID'].isin(good_sales['ID']))]
item_only = others[others['item_id'].isin(train['item_id'])]
no_data_items = others[~others['item_id'].isin(train['item_id'])]

print('1. Number of good pairs:', len(good_pairs))
print('2. Only Item_id Info:', len(item_only))

```

```
print('3. No Data Items:', len(no_data_items))
```

1. Number of good pairs: 111404
2. Only Item_id Info: 87550
3. No Data Items: 15246

This shows that, within the test set, there are 6,719 occurrences involving items that have not appeared previously in the training set, 5,615 occurrences where the shop-item combination has not occurred in the training set but the item has appeared with different shops, and 41,180 occurrences involving shop-item combinations that have been recorded in the training set.

1.1 Data Exploration

```
[4]: train.describe()
```

```
[4]:
```

	date_block_num	shop_id	item_id	item_price	item_cnt_day
count	2.935849e+06	2.935849e+06	2.935849e+06	2.935849e+06	2.935849e+06
mean	1.456991e+01	3.300173e+01	1.019723e+04	8.908532e+02	1.242641e+00
std	9.422988e+00	1.622697e+01	6.324297e+03	1.729800e+03	2.618834e+00
min	0.000000e+00	0.000000e+00	0.000000e+00	-1.000000e+00	-2.200000e+01
25%	7.000000e+00	2.200000e+01	4.476000e+03	2.490000e+02	1.000000e+00
50%	1.400000e+01	3.100000e+01	9.343000e+03	3.990000e+02	1.000000e+00
75%	2.300000e+01	4.700000e+01	1.568400e+04	9.990000e+02	1.000000e+00
max	3.300000e+01	5.900000e+01	2.216900e+04	3.079800e+05	2.169000e+03

1.1.1 Outliers and abnormal entries

```
[5]: negative_price = train[train.item_price < 0]
negative_price
```

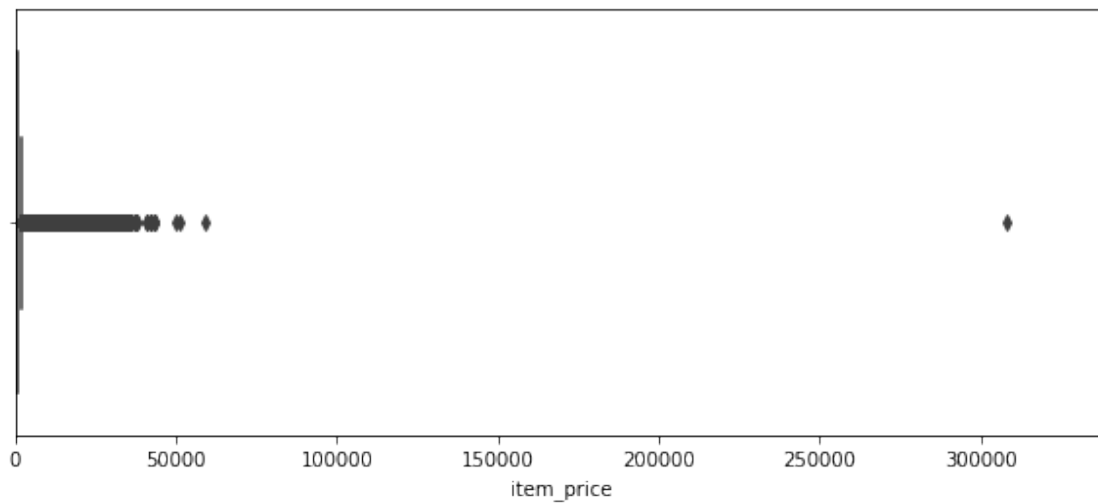
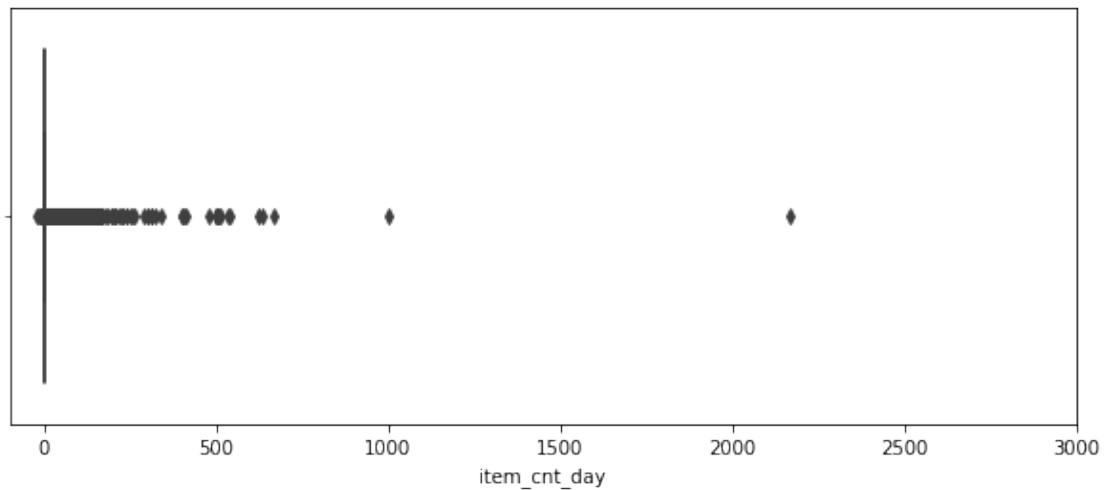
```
[5]:
```

	date	date_block_num	shop_id	item_id	item_price	item_cnt_day
484683	15.05.2013	4	32	2973	-1.0	1.0

```
[6]: #item count per day boxplot
plt.figure(figsize=(10,4))
plt.xlim(-100, 3000)
sns.boxplot(x=train.item_cnt_day)

#item price boxplot
plt.figure(figsize=(10,4))
plt.xlim(train.item_price.min(), train.item_price.max()*1.1)
sns.boxplot(x=train.item_price)
```

```
[6]: <matplotlib.axes._subplots.AxesSubplot at 0x7fa7b8cc8450>
```



These boxplots show outliers. To deal with these, entries with item price above 100000 and sales above 1001 will be removed from the dataset.

The data also includes negative entries for price (one item). This will be replaced with the median price.

```
[7]: train[train['item_id'] == 11373][['item_price']].sort_values(['item_price'])
```

```
[7]:      item_price
2909818    0.908714
2257993   38.500000
2048642   71.000000
1058343   72.200000
2462729   75.454545
```

```
...
2608320 1790.000000
885303 1821.000000
1006939 1931.000000
1830098 2111.500000
1058333 2137.000000
```

[464 rows x 1 columns]

```
[8]: train[train['item_id'] == 11365].sort_values(['item_price'])
```

```
[8]:
```

	date	date_block_num	shop_id	item_id	item_price \
1651714	16.05.2014	16	12	11365	124.0
2805487	21.08.2015	31	12	11365	170.0
1330776	13.01.2014	12	12	11365	180.0
1398688	25.02.2014	13	12	11365	194.0
661581	05.07.2013	6	12	11365	230.0
...
885161	28.09.2013	8	12	11365	9370.0
302568	12.03.2013	2	12	11365	10540.0
885165	23.09.2013	8	12	11365	11880.0
302544	05.03.2013	2	12	11365	14530.0
885138	17.09.2013	8	12	11365	59200.0

```

item_cnt_day
1651714      5.0
2805487      2.0
1330776      3.0
1398688      5.0
661581       4.0
...
885161       1.0
302568       1.0
885165       1.0
302544       1.0
885138       1.0
```

[242 rows x 6 columns]

```
[9]: #Correct train values
train['item_price'][2909818] = np.nan
train['item_cnt_day'][2909818] = np.nan
train['item_price'][2909818] = train[(train['shop_id'] ==12) &
↳(train['item_id'] == 11373) & (train['date_block_num'] == 33)]['item_price'].
↳median()
```

```

train['item_cnt_day'][2909818] = round(train[(train['shop_id'] ==12) &
↳(train['item_id'] == 11373) & (train['date_block_num'] ==
↳33)]['item_cnt_day'].median())
train['item_price'][885138] = np.nan
train['item_price'][885138] = train[(train['item_id'] == 11365) &
↳(train['shop_id'] ==12) & (train['date_block_num'] == 8)]['item_price'].
↳median()

```

```

[10]: #remove extreme outliers
train = train[train.item_price<100000]
train = train[train.item_cnt_day<1001]

```

```

[11]: #median of entries with shop id 32, item id 2973, month 4 and item price
↳greater than 0
median = train[(train.shop_id==32)&(train.item_id==2973)&(train.
↳date_block_num==4)&(train.item_price>0)].item_price.median()

#replace negative entry with median value
train.loc[train.item_price<0, 'item_price'] = median

```

```

[12]: #examine shop entries
shops

```

```

[12]:

```

		shop_name	shop_id
0	!	, 56	0
1	!	" "	1
2		" "	2
3	"	- "	3
4		" "	4
5		" "	5
6	(, 13)	6
7		" "	7
8	-	" "	8
9			9
10	.	39 ?	10
11	.	39 ^	11
12	-		12
13		" "	13
14	"	" II	14
15		"XXI "	15
16		" "	16
17	"	"	17
18		" "	18
19	"	"	19
20	"	"	20
21	"	"	21
22		21	22

23	"	"	(. 2)	23
24	"	"	(. 7)	24
25		"	"	25
26		"	" ()	26
27	"		II"	27
28	"		" II	28
29	"	"	()	29
30		"	"	30
31		"	"	31
32		"	"	32
33			"XL-3"	33
34		.	" "	34
35	.	"	"	35
36	"	"		36
37		"	"	37
38		"	"	38
39	"	"		39
40	"	"		40
41		"	"	41
42		"	"	42
43		"	"	43
44		"	"	44
45		"	"	45
46			"7 "	46
47		"	"	47
48	"	"		48
49		"	"	49
50		"	"	50
51	"	"		51
52		"	"	52
53		"	" 2	53
54		"	"	54
55		1 -		55
56		"	"	56
57			, 56	57
58	"	"		58
59	"	"		59

From looking at this dataframe, 4 duplicate entries have been found but listed separately due to spelling variations: + ! , 56 (shop id 0), , 56 (shop id 57) + !
 “ ” (shop id 1), “ ” (shop id 58)
 + . 39 ? (shop id 10), . 39 ² (shop id 11) + “ ”
 (shop id 39), “ ” (shop id 40)

These codes will be corrected to match in both the test and train sets.

From looking at the shop names, the first word is duplicated. From research, these are apparently the names of the cities in which these shops are located. This information will be useful for

processing the data later.

```
[13]: #           , 56
train.loc[train.shop_id == 0, 'shop_id'] = 57
test.loc[test.shop_id == 0, 'shop_id'] = 57

#           "           "
train.loc[train.shop_id == 1, 'shop_id'] = 58
test.loc[test.shop_id == 1, 'shop_id'] = 58

#           .           39 ²
train.loc[train.shop_id == 11, 'shop_id'] = 10
test.loc[test.shop_id == 11, 'shop_id'] = 10

#           "           "
train.loc[train.shop_id == 40, 'shop_id'] = 39
test.loc[test.shop_id == 40, 'shop_id'] = 39

#retain only shop_id present in test set
train = train.merge(test[['shop_id']].drop_duplicates(), how = 'inner')

#convert date to datetime
train['date'] = pd.to_datetime(train['date'], format = '%d.%m.%Y')
```

```
[14]: #save altered train
train.to_pickle('train_alt.pickle.gzde', compression='gzip')
```

```
[15]: del train
```

```
[16]: cats
```

```
[16]:
```

	item_category_name	item_category_id
0	PC - /	0
1	- PS2	1
2	- PS3	2
3	- PS4	3
4	- PSP	4
5	- PSVita	5
6	- XBOX 360	6
7	- XBOX ONE	7
8	()	8
9		9
10	- PS2	10
11	- PS3	11
12	- PS4	12
13	- PSP	13
14	- PSVita	14

15	- XBOX 360	15
16	- XBOX ONE	16
17	-	17
18	- PS2	18
19	- PS3	19
20	- PS4	20
21	- PSP	21
22	- PSVita	22
23	- XBOX 360	23
24	- XBOX ONE	24
25	-	25
26	Android -	26
27	MAC -	27
28	PC -	28
29	PC -	29
30	PC -	30
31	PC -	31
32	(, ,)	32
33	- Live!	33
34	- Live! ()	34
35	- PSN	35
36	- Windows ()	36
37	- Blu-Ray	37
38	- Blu-Ray 3D	38
39	- Blu-Ray 4K	39
40	- DVD	40
41	-	41
42	- ,	42
43	-	43
44	- ()	44
45	- 1	45
46	-	46
47	- ,	47
48	-	48
49	- 1	49
50	-	50
51	-	51
52	-	52
53	-	53
54	-	54
55	- CD	55
56	- CD	56
57	- MP3	57
58	-	58
59	-	59
60	-	60
61	-	61

```

62          - , , 62
63          - 63
64          - 64
65          - ( ) 65
66          - , 66
67          - 67
68          - , 68
69          - 69
70          - ( ) 70
71          - , , / 71
72          - 72
73          - 1 : 8 73
74          - MAC ( ) 74
75          - 75
76          - ( ) 76
77          - 77
78          - ( ) 78
79          79
80          - 80
81          ( ) 81
82          ( ) 82
83          83

```

```

[17]: # Select all duplicate rows based on one column
duplicateRowsCats = cats[cats.duplicated(['item_category_name'])]
print("Duplicate Rows based on a single column are:", duplicateRowsCats,
      ↪sep='\n')

```

Duplicate Rows based on a single column are:
Empty DataFrame
Columns: [item_category_name, item_category_id]
Index: []

No apparent issues have been spotted from the item category dataframe.

It has been noted that each category name includes type (duplicated) followed by subtype, separated by a dash.

```

[18]: items

```

```

[18]:
      item_name  item_id \
0          !      ( .)      D      0
1  !ABBY FineReader 12 Professional Edition Full...      1
2          ***      (UNV)      D      2
3          ***      (Univ)      D      3
4          ***      ( )      D      4
...          ...      ...
22165          2 [PC,      ]      22165

```

```

22166          1 :      [      ]      22166
22167          1 :      8 (+CD).    ...    22167
22168          Little Inu      22168
22169          (      )      22169

```

```

      item_category_id
0                40
1                76
2                40
3                40
4                40
...              ...
22165            31
22166            54
22167            49
22168            62
22169            69

```

[22170 rows x 3 columns]

```

[19]: # Select all duplicate rows based on one column
duplicateRowsItems = items[items.duplicated(['item_name'])]
print("Duplicate Rows based on a single column are:", duplicateRowsItems,
      ↪sep='\n')

```

Duplicate Rows based on a single column are:
Empty DataFrame
Columns: [item_name, item_id, item_category_id]
Index: []

No apparent issues have been spotted in the items dataframe.

1.1.2 Further exploration

```

[20]: train = pd.read_pickle('train_alt.pickle.gzde', compression='gzip')
train.nunique()

```

```

[20]: date                1034
date_block_num          34
shop_id                 42
item_id                21085
item_price              16567
item_cnt_day           191
dtype: int64

```

```

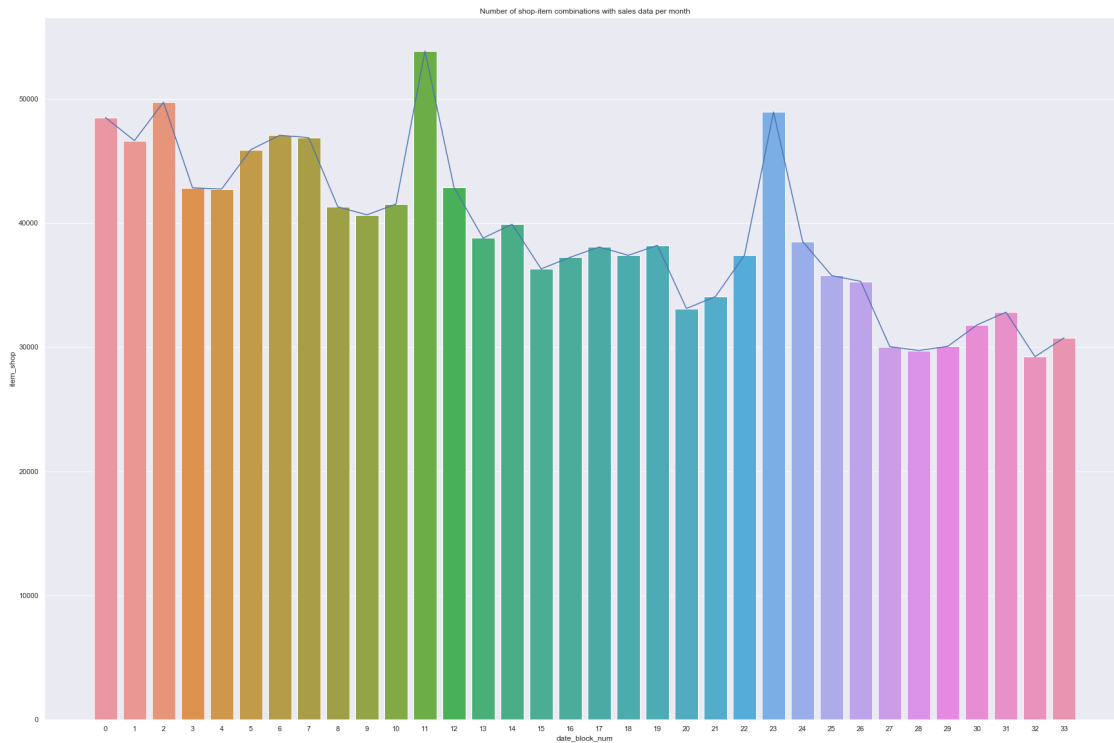
[21]: #number of sales-item combinations for which data is available per month
sns.set(rc={'figure.figsize':(30, 20)})
shop_item = pd.DataFrame(train[['date_block_num', 'shop_id',

```

```

                                'item_id']].drop_duplicates().
↳groupby('date_block_num').size().reset_index()
shop_item.columns = ['date_block_num', 'item_shop']
sns.barplot(x='date_block_num', y='item_shop', data=shop_item);
plt.plot(shop_item['item_shop']);
plt.title('Number of shop-item combinations with sales data per month')
del shop_item

```



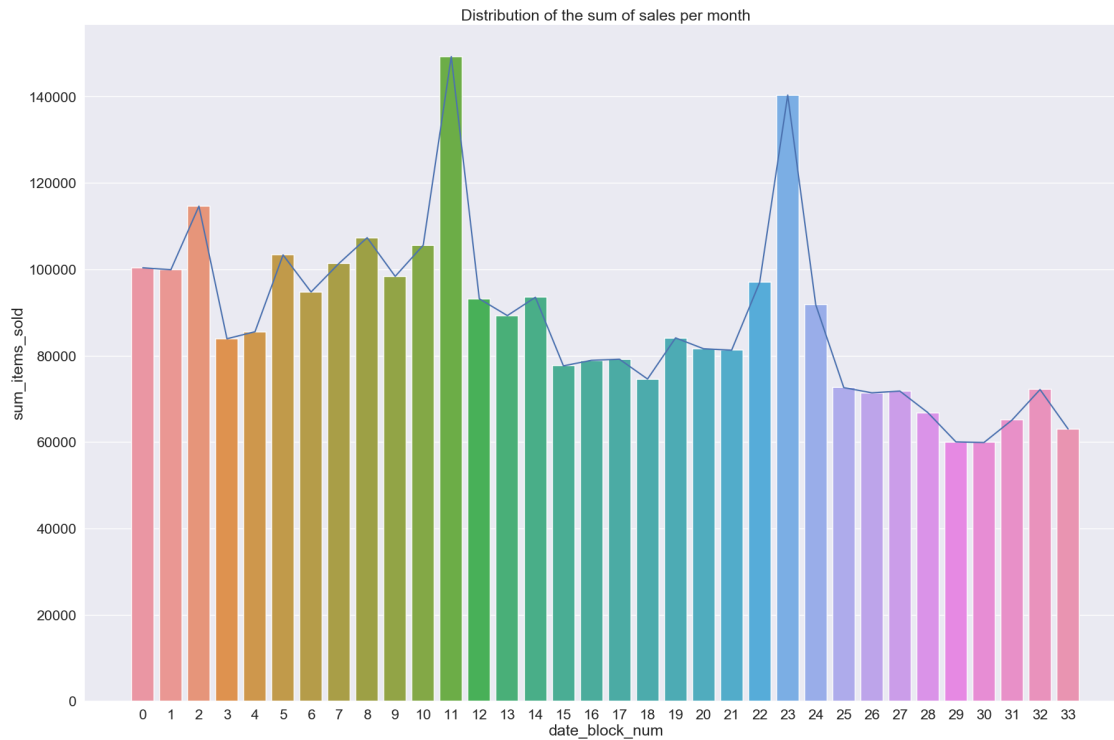
The plot above shows not all item-shop combinations have data for every month.

To enable calculation of these previously unrecorded combinations, the train dataset will need to be extended to include these missing combinations.

```

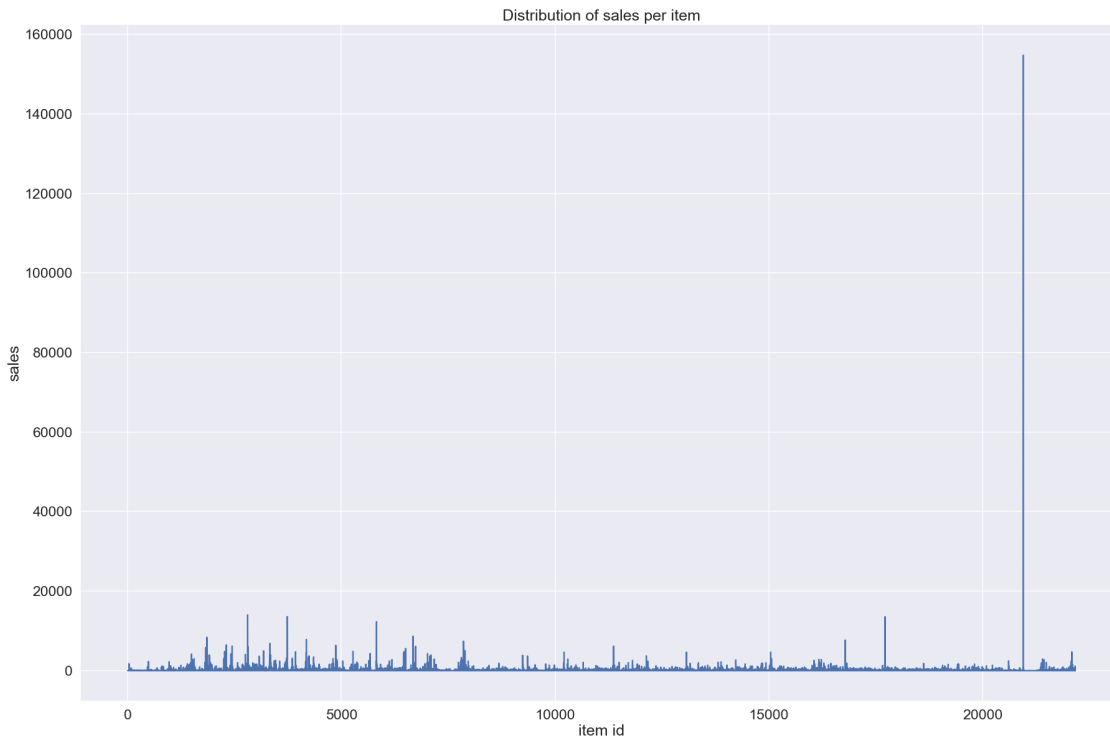
[22]: #visualise sales per month total
sns.set(rc={'figure.figsize':(30, 20)})
sns.set_context("talk", font_scale=1.4)
sales_month = pd.DataFrame(train.groupby(['date_block_num']).sum().
↳item_cnt_day).reset_index()
sales_month.columns = ['date_block_num', 'sum_items_sold']
sns.barplot(x='date_block_num', y='sum_items_sold',
            data=sales_month.reset_index());
plt.plot(sales_month.sum_items_sold)
plt.title('Distribution of the sum of sales per month')
del sales_month

```



This plot shows there is monthly variation in sales, therefore, date_block_num (month number) is a variable that should be accounted for.

```
[23]: #sales per item and shop combination
sns.set(rc={'figure.figsize':(30, 20)})
sns.set_context("talk", font_scale=1.4)
sales_item_id = pd.DataFrame(train.groupby(['item_id']).sum().item_cnt_day)
plt.xlabel('item id')
plt.ylabel('sales')
plt.title('Distribution of sales per item');
plt.plot(sales_item_id);
```



The plot shows that the vast majority of total sales for shop-item combinations are low, with some showing very high numbers.

```
[24]: sales_item_id = sales_item_id.reset_index()
      large_item = sales_item_id.item_cnt_day.argmax()
      large_item
```

```
[24]: 20055
```

```
[25]: sales_item_id.iloc[20055]
```

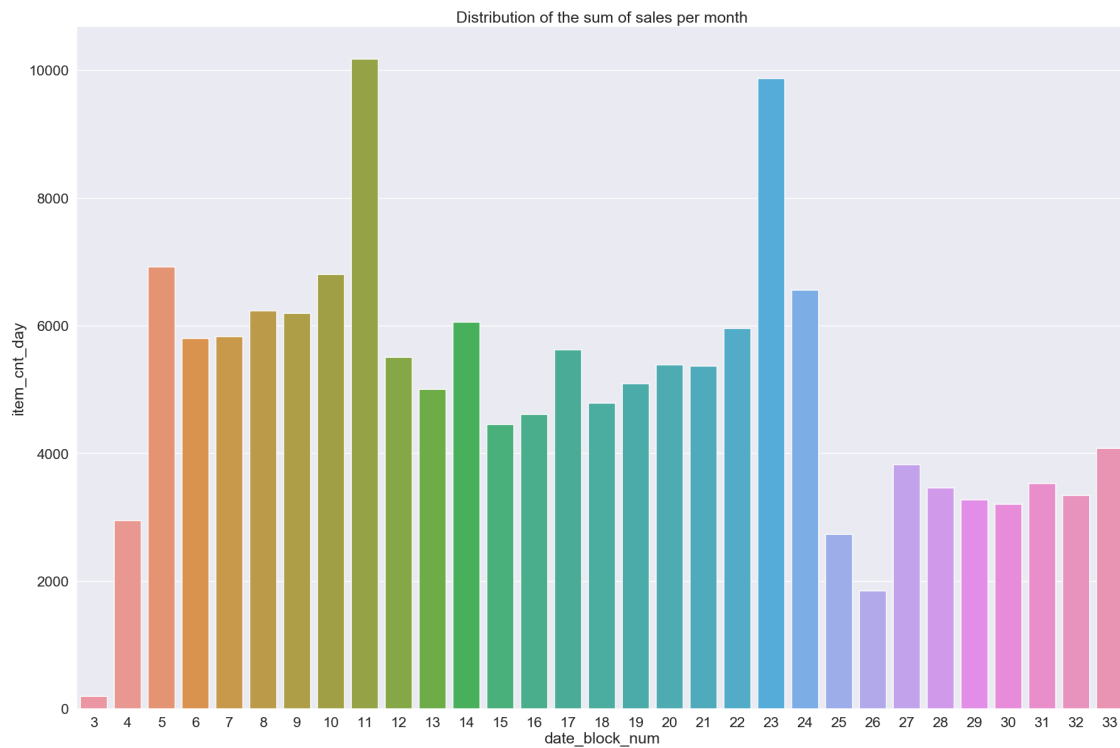
```
[25]: item_id      20949.0
      item_cnt_day 154771.0
      Name: 20055, dtype: float64
```

```
[26]: mask = (train.item_id == 20949)
      highest = train[mask]

      sns.set(rc={'figure.figsize':(30, 20)})
      sns.set_context("talk", font_scale=1.4)
      highest_df = pd.DataFrame(highest.groupby(['date_block_num']).sum().
      ↪ item_cnt_day)
      sns.barplot(x='date_block_num', y='item_cnt_day',
      data=highest_df.reset_index());
```

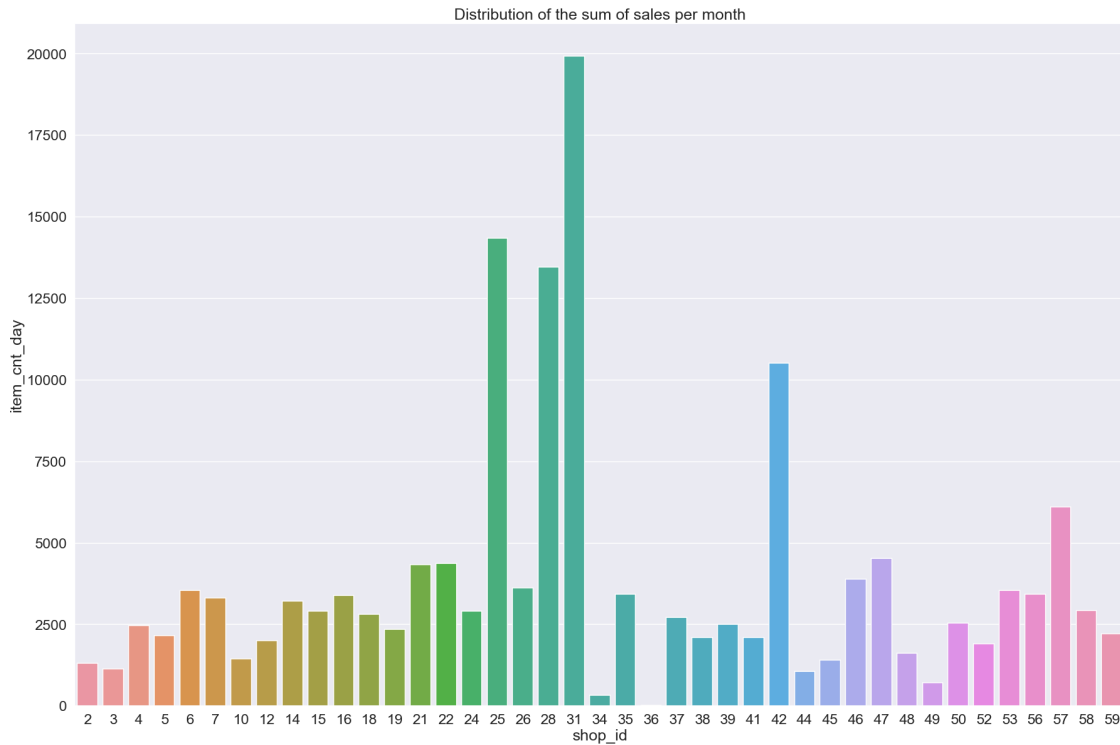
```
plt.title('Distribution of the sum of sales per month')
```

```
[26]: Text(0.5, 1.0, 'Distribution of the sum of sales per month')
```



```
[27]: sns.set(rc={'figure.figsize':(30, 20)})
sns.set_context("talk", font_scale=1.4)
highest_df2 = pd.DataFrame(highest.groupby(['shop_id']).sum().item_cnt_day)
sns.barplot(x='shop_id', y='item_cnt_day',
            data=highest_df2.reset_index());
plt.title('Distribution of the sum of sales per month')
```

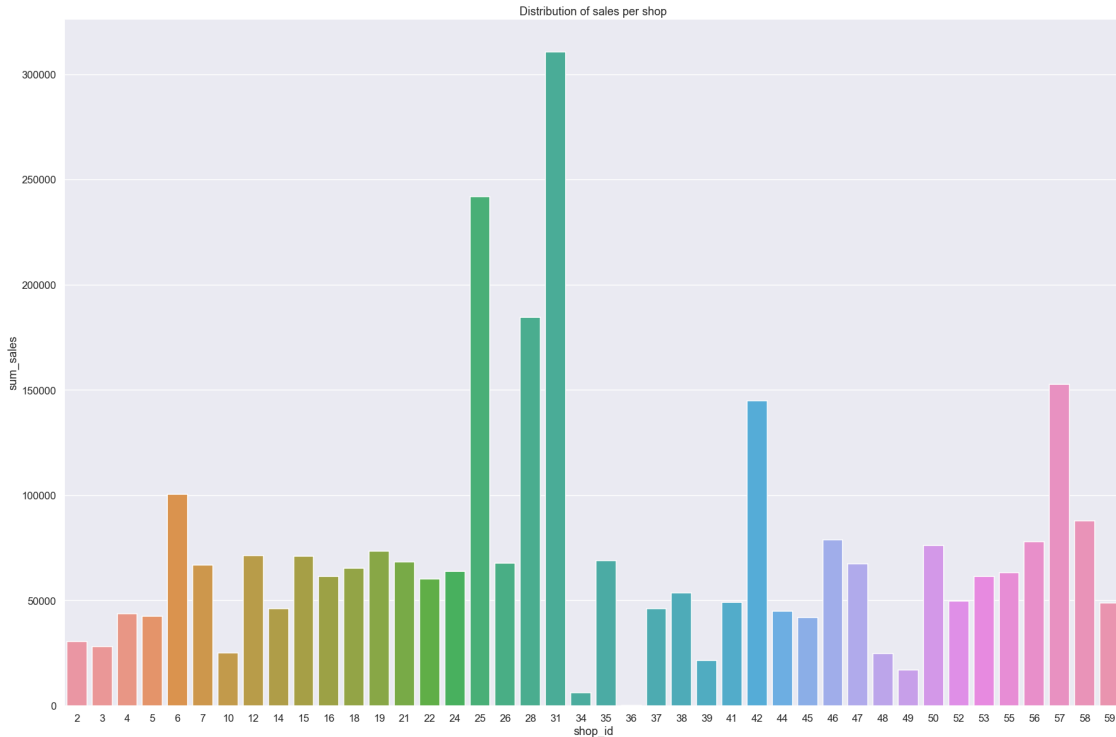
```
[27]: Text(0.5, 1.0, 'Distribution of the sum of sales per month')
```



The plots above show the distribution of sales for the item with the highest total sales within the dataset. The first shows some variation across months. The second shows a great deal of variation between shops for sales of this item, with a very small number of shops dominating and driving total sales of the item up.

Based on this analysis, total sales will be capped to prevent the possibility of extremely high sales for some items within certain stores driving up predictions to ranges beyond that which are likely.

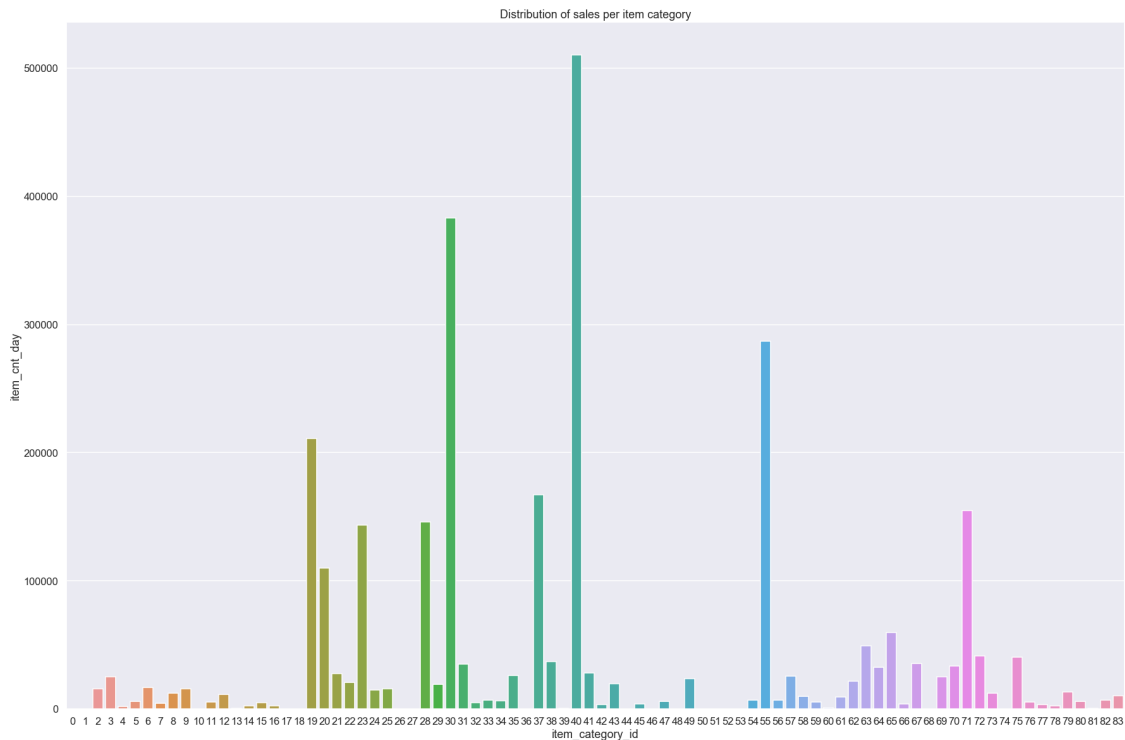
```
[28]: #sales per shop id
sns.set_context("talk", font_scale=1)
sales_month_shop_id = pd.DataFrame(train.groupby(['shop_id']).sum().
    ↪ item_cnt_day).reset_index()
sales_month_shop_id.columns = ['shop_id', 'sum_sales']
sns.barplot(x='shop_id', y='sum_sales', data=sales_month_shop_id)
plt.title('Distribution of sales per shop');
```

The above plot shows the distribution of sales of all items per shop. There is a lot of variation in sales per shop (possibly based on location or otherwise), and so this is a variable that the model will need to account for.

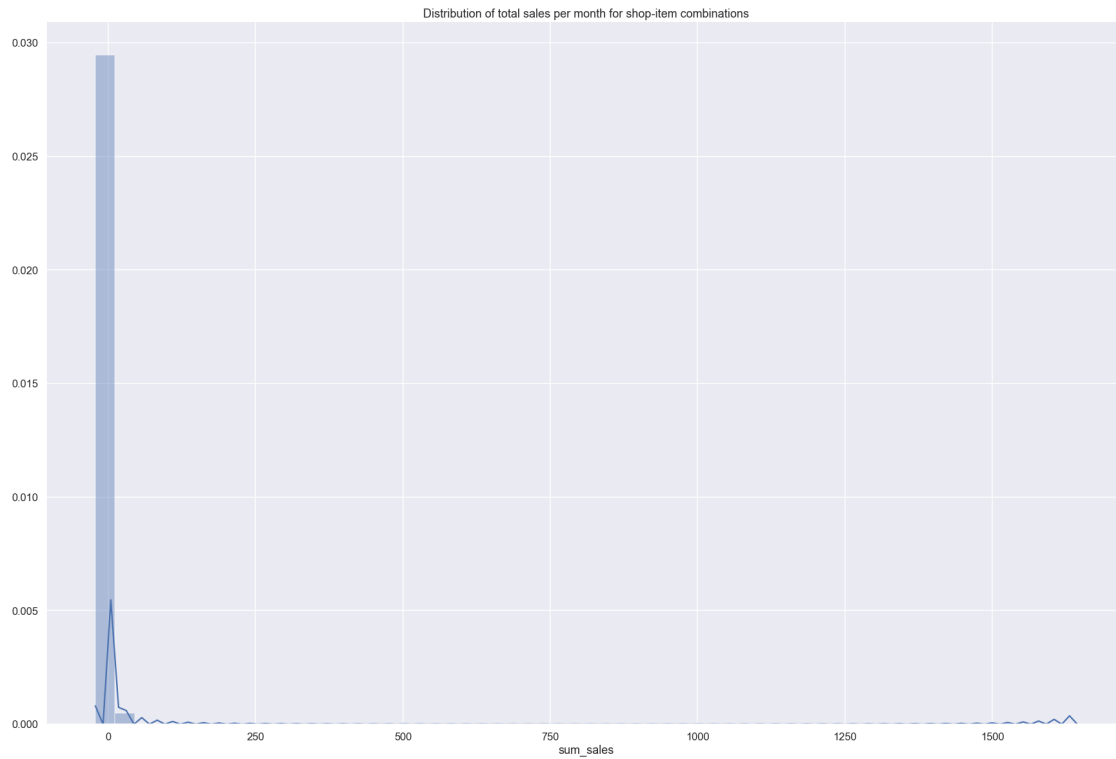
```
[29]: del sales_month_shop_id
      del sales_item_id
      del highest_df
      del highest
```

```
[30]: #sales per item category
sales_item_category = train.merge(items, how='left', on='item_id').
    ↳groupby('item_category_id')['item_cnt_day'].sum()
sns.barplot(x = 'item_category_id', y = 'item_cnt_day', data =
    ↳sales_item_category.reset_index());
plt.title('Distribution of sales per item category');
del sales_item_category
```

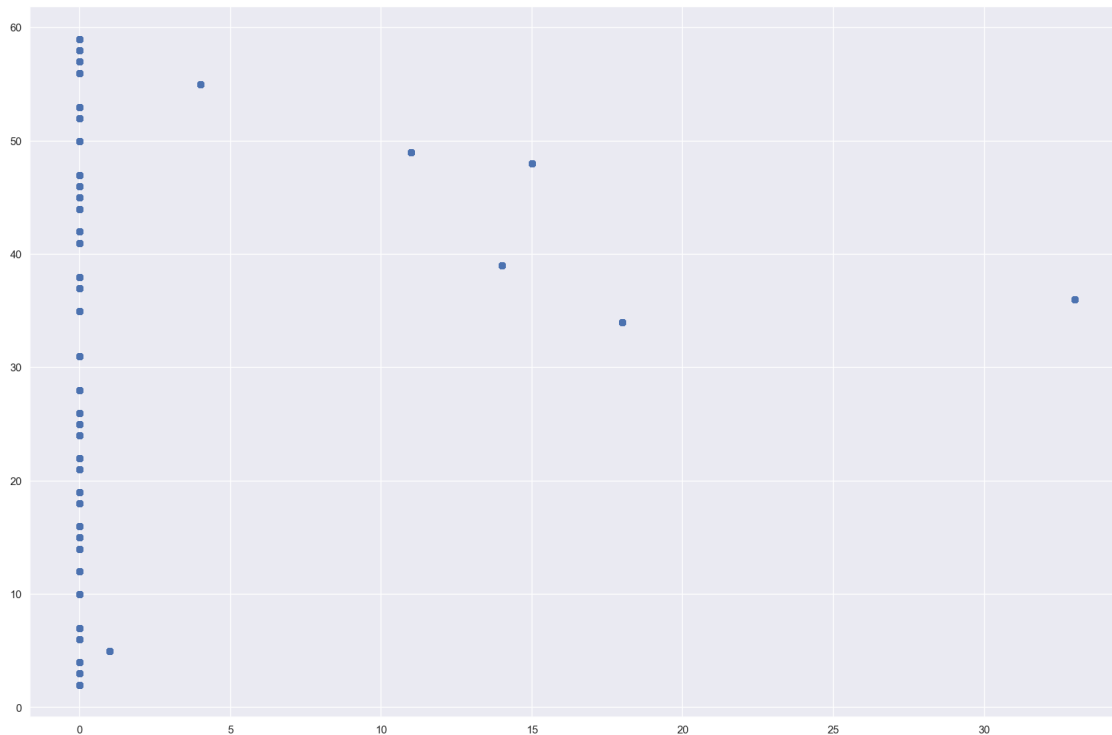


Again, there is a lot of variation in sales per item category id, which will likely be helpful in predicting sales.

```
[31]: #sales per shop id
sns.set_context("talk", font_scale=1)
sales_month_shop_item = pd.DataFrame(train.
    ↳groupby(['date_block_num', 'shop_id', 'item_id']).sum().item_cnt_day).
    ↳reset_index()
sales_month_shop_item.columns =
    ↳['date_block_num', 'shop_id', 'item_id', 'sum_sales']
sns.distplot(sales_month_shop_item['sum_sales'])
plt.title('Distribution of total sales per month for shop-item combinations');
del sales_month_shop_item
```



```
[32]: #scatter plot of item_id and date_block_num
sales1 = train.copy()
sales1['min_month'] = sales1.groupby(['shop_id'])['date_block_num'].
    ↪transform('min')
plt.scatter(sales1.min_month, sales1.shop_id)
del sales1
```



```
[33]: #scatter plot of item_id and date_block_num  
sales1 = train.copy()  
sales1['min_month'] = sales1.groupby(['item_id'])['date_block_num'].  
    ↪transform('min')  
plt.figure(figsize=(20,40))  
plt.scatter(sales1.min_month, sales1.item_id)  
del sales1
```



There is no apparent pattern for item id numbers overall - it does not appear to be the case that the later an item is released, the larger the id number.

This is a similar story for shop_id; no apparent relationship exists between shop_id and sales.

1.2 Data Preprocessing

1.2.1 Shops dataframe

This involves feature extraction from text, namely splitting city name from shop name and label encoded. These codes will be used later for mean encoding.

```
[34]: shops['shop_name'] = shops['shop_name'].apply(lambda x: x.lower()).str.  
      ↪replace('[^\w\s]', '').str.replace('\d+', '').str.strip()  
shops['shop_city'] = shops['shop_name'].str.partition(' ')[0]  
shops['shop_type'] = shops['shop_name'].apply(lambda x: ' ' if ' ' in x else  
      ↪' ' if ' ' in x else ' ' if ' ' in x else ' ' if ' ' in x else ' ' if '  
      ↪in x else 'NO_DATA')  
shops.head()
```

```
[34]:
```

	shop_name	shop_id	shop_city	shop_type
0		0		NO_DATA
1		1		
2		2		
3		3		
4		4		

```
[35]: #label encode city name  
shops['city_id'] = LabelEncoder().fit_transform(shops['shop_city'])  
shops['shop_type_id'] = LabelEncoder().fit_transform(shops['shop_type'])  
  
#retain only numerical codes  
shops = shops[['shop_id', 'shop_type_id', 'city_id']]  
  
shops.head()
```

```
[35]:
```

	shop_id	shop_type_id	city_id
0	0	0	29
1	1	5	29
2	2	5	0
3	3	3	1
4	4	5	2

1.2.2 Category dataframe

As noted in exploration, category name contains type and subtype of item. This section involves feature extraction from text, namely splitting item category name into type and subtype of the

item and label encoding. These codes will be used later for the purposes of mean encoding.

```
[36]: #split and separate out item type
cats['split'] = cats['item_category_name'].str.split('-')
cats['type'] = cats['split'].map(lambda x: x[0].strip())
cats.head()
```

```
[36]:      item_category_name  item_category_id      split \
0  PC - / 0 [PC , / ]
1      - PS2 1 [ , PS2]
2      - PS3 2 [ , PS3]
3      - PS4 3 [ , PS4]
4      - PSP 4 [ , PSP]

      type
0      PC
1
2
3
4
```

```
[37]: #label encode item type
cats['cat_type_id'] = LabelEncoder().fit_transform(cats['type'])
```

```
[38]: # separate out item subtype, if no subtype replace with type
cats['subtype'] = cats['split'].map(lambda x: x[1].strip() if len(x) > 1 else
    ↪x[0].strip())

#label encode item subtype
cats['cat_subtype_id'] = LabelEncoder().fit_transform(cats['subtype'])
```

```
[39]: #retain only nnumerical codes
cats = cats[['item_category_id', 'cat_type_id', 'cat_subtype_id']]
cats.head()
```

```
[39]:      item_category_id  cat_type_id  cat_subtype_id
0              0          0          29
1              1          1           9
2              2          1          10
3              3          1          11
4              4          1          13
```

1.2.3 Items dataframe

```
[259]: items.head()
```

```
[259]:
```

		item_name	item_id	\
0	!	(.)	D	0
1	!ABBY FineReader 12 Professional Edition Full...			1
2	***	(UNV)	D	2
3	***	(Univ)	D	3
4	***	()	D	4

	item_category_id
0	40
1	76
2	40
3	40
4	40

Encode “features” that many items have.

The structure is always the same Item name [category feature] (additional feature)

This can be split, and encoded.

```
[260]: from collections import Counter
from operator import itemgetter
items['name_1'], items['name_2'] = items['item_name'].str.split('[', 1).str
items['name_1'], items['name_3'] = items['item_name'].str.split('(', 1).str

items['name_2'] = items['name_2'].str.replace('[^A-Za-z0-9 - -]+', ' ').str.
    ↪lower()
items['name_3'] = items['name_3'].str.replace('[^A-Za-z0-9 - -]+', ' ').str.
    ↪lower()
items = items.fillna('0')

result_1 = Counter(' '.join(items['name_2'].values.tolist()).split(' ')).items()
result_1 = sorted(result_1, key=itemgetter(1))
result_1 = pd.DataFrame(result_1, columns=['feature', 'count'])
result_1 = result_1[(result_1['feature'].str.len() > 1) & (result_1['count'] >_
    ↪200)]

result_2 = Counter(' '.join(items['name_3'].values.tolist()).split(" ")).items()
result_2 = sorted(result_2, key=itemgetter(1))
result_2 = pd.DataFrame(result_2, columns=['feature', 'count'])
result_2 = result_2[(result_2['feature'].str.len() > 1) & (result_2['count'] >_
    ↪200)]

result = pd.concat([result_1, result_2])
result = result.drop_duplicates(subset=['feature'])

print('Most common additional features:', result)
```


	Most common additional features:	feature	count
130		284	
131		340	
132		399	
133		400	
134		360	465
135	jewel		552
136	xbox		589
137	ps3		611
138		1428	
139		1995	
140	pc		2585
141		3427	
1981	box		246
1983	3d		409
1985	dvd		503
1986	digipack		541
1988		757	
1991	mp3		854
1992	cd		871
1993		1849	
1994	bd		2320

Item name correction

For our basic “name feature” it is enough to find identical items (not similar but identical),

```
[261]: print('Unique item names:', len(items['item_name'].unique()))
```

Unique item names: 22170

```
[262]: items.name_1.nunique(), items.name_2.nunique(), items.name_3.nunique(), items.
       ↪ item_category_id.nunique()
```

```
[262]: (20611, 175, 1666, 84)
```

```
[263]: import re
def name_correction(x):
    x = x.lower()
    x = x.partition('[')[0]
    x = x.partition('(')[0]
    x = re.sub('[^A-Za-z0-9- -]+', ' ', x)
    x = x.replace(' ', ' ')
    x = x.strip()
    return x

items['name_1'] = items['name_1'].apply(lambda x: name_correction(x))
items.head()
```

```
[263]:
```

		item_name	item_id	\
0	!	(.)	D	0
1	!ABBY	FineReader 12 Professional Edition Full...		1
2	***	(UNV)	D	2
3	***	(Univ)	D	3
4	***	()	D	4

	item_category_id	name_1	\
0	40		
1	76	abby finereader 12 professional edition full	
2	40		
3	40		
4	40		

	name_2	name_3
0	0	d
1	pc	0
2	0	unv d
3	0	univ d
4	0	d

```
[264]: items.name_1.nunique(), items.name_2.nunique(), items.name_3.nunique(), items.
      ↪ item_category_id.nunique()
```

```
[264]: (18121, 175, 1666, 84)
```

```
[45]: print('Unique item names after correction:', len(items['item_name'].unique()))
```

Unique item names after correction: 18121

```
[149]: #label encode name_1, 2, and 3
      #items['item_name_id'] = LabelEncoder().fit_transform(items['name_1'])

      items['item_type_id'] = LabelEncoder().fit_transform(items['name_2'])
      items['item_subtype_id'] = LabelEncoder().fit_transform(items['name_3'])

      items = items.drop(['item_name', 'name_2', 'name_3'], axis=1)
```

1.2.4 Monthly sales

As noted previously, some item-shop combinations in the test set are not present in the train set - the train set includes only items that have been sold in the past, thus the test items not present in the train have not, and as predictions are for the purposes of this task in the future, the target value should be zero. To ensure these items are accounted for in the model, all possible item-shop combinations need to be included for each month and set at zero.

Furthermore, the task is to predict monthly sales. So, daily sales reported need to be aggregated to monthly.

```
[48]: # For every month create a grid from all shops/items combinations from that
      ↪ month

grid = []

for block_num in train['date_block_num'].unique():
    cur_shops = train[train['date_block_num']==block_num]['shop_id'].unique()
    cur_items = train[train['date_block_num']==block_num]['item_id'].unique()
    grid.append(np.array(list(product(*[cur_shops, cur_items,
    ↪ [block_num]])), dtype='int32'))

# #turn the grid into pandas dataframe
index_cols = ['shop_id', 'item_id', 'date_block_num']
grid = pd.DataFrame(np.vstack(grid), columns = index_cols, dtype=np.int32)

index_cols = ['shop_id', 'item_id', 'date_block_num']
train['item_cnt_day'] = train['item_cnt_day'].clip(0,20)
gb_cnt = train.groupby(index_cols)['item_cnt_day'].agg(['sum']).reset_index().
    ↪ rename(columns = {'sum': 'item_cnt_month'})
gb_cnt['item_cnt_month'] = gb_cnt['item_cnt_month'].clip(0,20).astype(np.int)

[49]: #join aggregated data to the grid
train = pd.merge(grid, gb_cnt, how='left', on=index_cols).fillna(0)
train['item_cnt_month'] = train['item_cnt_month'].astype(int)
train = downcast_dtypes(train)

[50]: #sort the data
train.sort_values(['date_block_num', 'shop_id', 'item_id'], inplace=True)

[51]: #add additional column for shop-item
train['shop_item_id'] = train['shop_id'].apply(str) + '_' + train['item_id'].
    ↪ apply(str)
test['shop_item_id'] = test['shop_id'].apply(str) + '_' + test['item_id'].
    ↪ apply(str)
```

1.2.5 Add codes from category, item and shop dataframes to main dataset

These codes are needed for mean encoding next.

```
[52]: sales = pd.merge(train, shops, on=['shop_id'], how='left')
sales = pd.merge(sales, items, on=['item_id'], how='left')
sales = pd.merge(sales, cats, on=['item_category_id'], how='left')

test = pd.merge(test, shops, on=['shop_id'], how='left')
test = pd.merge(test, items, on=['item_id'], how='left')
test = pd.merge(test, cats, on=['item_category_id'], how='left')
```

```
[53]: #reduce data size
sales = downcast_dtypes(sales)
```

```
[54]: del shops
del items
del cats
```

```
[55]: sales.head()
```

```
[55]:
```

	shop_id	item_id	date_block_num	item_cnt_month	shop_item_id	\
0	2	19	0	0	2_19	
1	2	27	0	1	2_27	
2	2	28	0	0	2_28	
3	2	29	0	0	2_29	
4	2	32	0	0	2_32	

	shop_type_id	city_id	item_category_id	item_type_id	item_subtype_id	\
0	5	0	40	4	42	
1	5	0	19	77	42	
2	5	0	30	108	42	
3	5	0	23	124	42	
4	5	0	40	4	42	

	cat_type_id	cat_subtype_id
0	11	4
1	5	10
2	8	55
3	5	16
4	11	4

1.2.6 Mean encode features

Mean encode categorical features using KFold, LOO, Smoothing and Expanding and select the version of each feature with the highest correlation coefficient.

As item_item_month for the test set will not be available to include in mean encoding, these encodings will be conducted on the data minus the test set.

```
[56]: #global mean set at group mean as train_df produces NaN due to large number of
      ↳ 0 values
mean_encoded_col =
      ↳ ['shop_id', 'item_id', 'shop_item_id', 'shop_type_id', 'city_id', 'item_category_id', 'item_type_
        ↳ item_subtype_id', 'cat_type_id', 'cat_subtype_id']

from tqdm import tqdm
from sklearn.model_selection import KFold

Target = 'item_cnt_month'
```

```

global_mean = sales[Target].mean()
y_tr = sales[Target].values

for col in tqdm(mean_encoded_col):
    col_tr = sales[[col] + [Target]]
    corrcoeffs = pd.DataFrame(columns = ['Cor'])

    # 3.1.1 Mean encodings - KFold scheme
    kf = KFold(n_splits = 5, shuffle = False, random_state = 0)
    col_tr[col + '_cnt_month_mean_Kfold'] = np.nan

    for tr_ind, val_ind in kf.split(col_tr):
        X_tr, X_val = col_tr.iloc[tr_ind], col_tr.iloc[val_ind]
        means = X_val[col].map(X_tr.groupby(col)[Target].mean())
        X_val[col + '_cnt_month_mean_Kfold'] = means
        col_tr.iloc[val_ind] = X_val

    col_tr.fillna(global_mean, inplace = True)
    corrcoeffs.loc[col + '_cnt_month_mean_Kfold'] = np.corrcoef(y_tr, col_tr[col +
    ↪ '_cnt_month_mean_Kfold'])[0][1]

    # 3.1.2 Mean encodings - Leave-one-out scheme
    item_id_target_sum = col_tr.groupby(col)[Target].sum()
    item_id_target_count = col_tr.groupby(col)[Target].count()
    col_tr[col + '_cnt_month_sum'] = col_tr[col].map(item_id_target_sum)
    col_tr[col + '_cnt_month_count'] = col_tr[col].map(item_id_target_count)
    col_tr[col + '_target_mean_L00'] = (col_tr[col + '_cnt_month_sum'] -
    ↪ col_tr[Target]) / (col_tr[col + '_cnt_month_count'] - 1)
    col_tr.fillna(global_mean, inplace = True)
    corrcoeffs.loc[col + '_target_mean_L00'] = np.corrcoef(y_tr, col_tr[col +
    ↪ '_target_mean_L00'])[0][1]

    # 3.1.3 Mean encodings - Smoothing
    item_id_target_mean = col_tr.groupby(col)[Target].mean()
    item_id_target_count = col_tr.groupby(col)[Target].count()
    col_tr[col + '_cnt_month_mean'] = col_tr[col].map(item_id_target_mean)
    col_tr[col + '_cnt_month_count'] = col_tr[col].map(item_id_target_count)
    alpha = 100
    col_tr[col + '_cnt_month_mean_Smooth'] = (col_tr[col + '_cnt_month_mean'] *
    ↪ col_tr[col + '_cnt_month_count'] + global_mean * alpha) / (alpha +
    ↪ col_tr[col + '_cnt_month_count'])
    col_tr[col + '_cnt_month_mean_Smooth'].fillna(global_mean, inplace=True)
    corrcoeffs.loc[col + '_cnt_month_mean_Smooth'] = np.corrcoef(y_tr,
    ↪ col_tr[col + '_cnt_month_mean_Smooth'])[0][1]

```

```

# 3.1.4 Mean encodings - Expanding mean scheme
cumsum = col_tr.groupby(col)[Target].cumsum() - col_tr[Target]
sumcnt = col_tr.groupby(col).cumcount()
col_tr[col + '_cnt_month_mean_Expanding'] = cumsum / sumcnt
col_tr[col + '_cnt_month_mean_Expanding'].fillna(global_mean, inplace=True)
corrcoefs.loc[col + '_cnt_month_mean_Expanding'] = np.corrcoef(y_tr,
→col_tr[col + '_cnt_month_mean_Expanding'])[0][1]

sales = pd.concat([sales, col_tr[corrcoefs['Cor'].idxmax()]], axis = 1)
print(corrcoefs.sort_values('Cor'))

```

10%| | 1/10 [00:04<00:38, 4.28s/it]

	Cor
shop_id_cnt_month_mean_Kfold	0.172836
shop_id_target_mean_L00	0.174991
shop_id_cnt_month_mean_Smooth	0.175016
shop_id_cnt_month_mean_Expanding	0.175150

20%| | 2/10 [00:09<00:36, 4.53s/it]

	Cor
item_id_cnt_month_mean_Kfold	0.312957
item_id_cnt_month_mean_Smooth	0.479641
item_id_target_mean_L00	0.481724
item_id_cnt_month_mean_Expanding	0.565665

30%| | 3/10 [00:51<01:50, 15.81s/it]

	Cor
shop_item_id_cnt_month_mean_Kfold	0.423936
shop_item_id_cnt_month_mean_Expanding	0.542637
shop_item_id_target_mean_L00	0.577498
shop_item_id_cnt_month_mean_Smooth	0.600000

40%| | 4/10 [00:55<01:13, 12.30s/it]

	Cor
shop_type_id_cnt_month_mean_Kfold	0.034009
shop_type_id_target_mean_L00	0.037720
shop_type_id_cnt_month_mean_Smooth	0.037738
shop_type_id_cnt_month_mean_Expanding	0.039478

50%| | 5/10 [00:59<00:49, 9.83s/it]

	Cor
city_id_cnt_month_mean_Kfold	0.117109
city_id_target_mean_L00	0.119897
city_id_cnt_month_mean_Smooth	0.119918
city_id_cnt_month_mean_Expanding	0.120103

60%| | 6/10 [01:04<00:32, 8.21s/it]

Cor

item_category_id_cnt_month_mean_Kfold	0.270138
item_category_id_cnt_month_mean_Smooth	0.289958
item_category_id_target_mean_L00	0.289988
item_category_id_cnt_month_mean_Expanding	0.292890

70%| | 7/10 [01:08<00:21, 7.14s/it]

Cor

item_type_id_cnt_month_mean_Kfold	0.194394
item_type_id_cnt_month_mean_Smooth	0.213486
item_type_id_target_mean_L00	0.213759
item_type_id_cnt_month_mean_Expanding	0.225101

80%| | 8/10 [01:13<00:12, 6.33s/it]

Cor

item_subtype_id_cnt_month_mean_Kfold	0.221293
item_subtype_id_cnt_month_mean_Smooth	0.250972
item_subtype_id_target_mean_L00	0.251816
item_subtype_id_cnt_month_mean_Expanding	0.258824

90%| | 9/10 [01:17<00:05, 5.71s/it]

Cor

cat_type_id_cnt_month_mean_Kfold	0.155011
cat_type_id_target_mean_L00	0.169346
cat_type_id_cnt_month_mean_Smooth	0.169398
cat_type_id_cnt_month_mean_Expanding	0.174876

100%| | 10/10 [01:21<00:00, 8.19s/it]

Cor

cat_subtype_id_cnt_month_mean_Kfold	0.270167
cat_subtype_id_cnt_month_mean_Smooth	0.288706
cat_subtype_id_target_mean_L00	0.288749
cat_subtype_id_cnt_month_mean_Expanding	0.290632

1.2.7 Combine test and train

```
[57]: if Validation == False:
      test['date_block_num'] = 34
      all_data = pd.concat([sales, test], axis = 0)
      all_data = all_data.drop(columns = ['ID'])

else:
    all_data = sales
```

1.2.8 Feature Generation

1.2.9 1. Create time lagged features

```
[58]: sales = downcast_dtypes(all_data)
```

```
[59]: #function to create time lags
def lag_feature(df, lags, col):
    tmp = df[['date_block_num', 'shop_id', 'item_id', col]]
    for i in lags:
        shifted = tmp.copy()
        shifted.columns = ['date_block_num', 'shop_id', 'item_id',
        ↪col+' _lag_'+str(i)]
        shifted['date_block_num'] += i
        df = pd.merge(df, shifted, on=['date_block_num', 'shop_id', 'item_id'],
        ↪how='left')
    return df
```

```
[60]: #lagging sales for shop-item combinations per month
sales = lag_feature(sales, [1,2,3,4,5,6,7,8,9,10,11,12], 'item_cnt_month')

#as prediction months will have no mean encoded features and some of these vary
    ↪by month,
#they will be lagged with current ones removed

cols = ['shop_id_cnt_month_mean_Expanding',
        'item_id_cnt_month_mean_Expanding',
        'shop_item_id_cnt_month_mean_Smooth',
        'shop_type_id_cnt_month_mean_Expanding',
        'city_id_cnt_month_mean_Expanding',
        'item_category_id_cnt_month_mean_Expanding',
        'item_type_id_cnt_month_mean_Expanding',
        'item_subtype_id_cnt_month_mean_Expanding',
        'cat_type_id_cnt_month_mean_Expanding',
        'cat_subtype_id_cnt_month_mean_Expanding']

shift_range = [1, 2, 3, 4, 12]

for col in cols:
    sales = lag_feature(sales, shift_range, col)

sales.head()
```

```
[60]:   cat_subtype_id  cat_subtype_id_cnt_month_mean_Expanding  cat_type_id \
0                4                                0.311493         11
1               10                                0.311493          5
2               55                                0.311493          8
3               16                                0.311493          5
```


4	4	0.000000	11
---	---	----------	----

	cat_type_id_cnt_month_mean_Expanding	city_id \
0	0.311493	0
1	0.311493	0
2	0.311493	0
3	1.000000	0
4	0.000000	0

	city_id_cnt_month_mean_Expanding	date_block_num	item_category_id \
0	0.311493	0	40
1	0.000000	0	19
2	0.500000	0	30
3	0.333333	0	23
4	0.250000	0	40

	item_category_id_cnt_month_mean_Expanding	item_cnt_month	item_id \
0	0.311493	0.0	19
1	0.311493	1.0	27
2	0.311493	0.0	28
3	0.311493	0.0	29
4	0.000000	0.0	32

	item_id_cnt_month_mean_Expanding	item_subtype_id \
0	0.311493	42
1	0.311493	42
2	0.311493	42
3	0.311493	42
4	0.311493	42

	item_subtype_id_cnt_month_mean_Expanding	item_type_id \
0	0.311493	4
1	0.000000	77
2	0.500000	108
3	0.333333	124
4	0.250000	4

	item_type_id_cnt_month_mean_Expanding	shop_id \
0	0.311493	2
1	0.311493	2
2	0.311493	2
3	0.311493	2
4	0.000000	2

	shop_id_cnt_month_mean_Expanding	shop_item_id \
0	0.311493	2_19
1	0.000000	2_27

2	0.500000	2_28
3	0.333333	2_29
4	0.250000	2_32

	shop_item_id_cnt_month_mean_Smooth	shop_type_id \
0	0.308409	5
1	0.288254	5
2	0.275657	5
3	0.291115	5
4	0.314547	5

	shop_type_id_cnt_month_mean_Expanding	item_cnt_month_lag_1 \
0	0.311493	NaN
1	0.000000	NaN
2	0.500000	NaN
3	0.333333	NaN
4	0.250000	NaN

	item_cnt_month_lag_2	item_cnt_month_lag_3	...	\
0	NaN	NaN	...	
1	NaN	NaN	...	
2	NaN	NaN	...	
3	NaN	NaN	...	
4	NaN	NaN	...	

	item_category_id_cnt_month_mean_Expanding_lag_1 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_category_id_cnt_month_mean_Expanding_lag_2 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_category_id_cnt_month_mean_Expanding_lag_3 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_category_id_cnt_month_mean_Expanding_lag_4 \
--	---

0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_category_id_cnt_month_mean_Expanding_lag_12 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_type_id_cnt_month_mean_Expanding_lag_1 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_type_id_cnt_month_mean_Expanding_lag_2 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_type_id_cnt_month_mean_Expanding_lag_3 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_type_id_cnt_month_mean_Expanding_lag_4 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_type_id_cnt_month_mean_Expanding_lag_12 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_subtype_id_cnt_month_mean_Expanding_lag_1	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	
4	NaN	

	item_subtype_id_cnt_month_mean_Expanding_lag_2	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	
4	NaN	

	item_subtype_id_cnt_month_mean_Expanding_lag_3	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	
4	NaN	

	item_subtype_id_cnt_month_mean_Expanding_lag_4	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	
4	NaN	

	item_subtype_id_cnt_month_mean_Expanding_lag_12	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	
4	NaN	

	cat_type_id_cnt_month_mean_Expanding_lag_1	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	
4	NaN	

	cat_type_id_cnt_month_mean_Expanding_lag_2	\
0	NaN	
1	NaN	
2	NaN	

3	NaN
4	NaN

	cat_type_id_cnt_month_mean_Expanding_lag_3 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_type_id_cnt_month_mean_Expanding_lag_4 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_type_id_cnt_month_mean_Expanding_lag_12 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_subtype_id_cnt_month_mean_Expanding_lag_1 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_subtype_id_cnt_month_mean_Expanding_lag_2 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_subtype_id_cnt_month_mean_Expanding_lag_3 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_subtype_id_cnt_month_mean_Expanding_lag_4 \
0	NaN

1	NaN
2	NaN
3	NaN
4	NaN

cat_subtype_id_cnt_month_mean_Expanding_lag_12	
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

[5 rows x 84 columns]

```
[61]: #as these will not be present for test set, only the lagged ones will be
      ↪retained
sales = sales.drop(cols, axis=1)
```

```
[62]: sales = downcast_dtypes(sales)
```

```
[63]: sales.to_pickle('data_1.pickle.gzde', compression='gzip')
```

```
[64]: del sales
```

1.2.10 2. Item price trend over the previous 6 months

Price will influence demand for a product, and thus sales.

```
[65]: sales = pd.read_pickle('data_1.pickle.gzde', compression='gzip')
train = pd.read_pickle('train_alt.pickle.gzde', compression='gzip')
```

```
[66]: # calculate mean item price
group = train.groupby(['item_id']).agg({'item_price': ['mean']})
group.columns = ['avg_item_price']
group.reset_index(inplace=True)

#add to df
sales = pd.merge(sales, group, on=['item_id'], how='left')
sales['avg_item_price'] = sales['avg_item_price'].astype(np.float16)
```

```
[67]: # calculate mean item price per month
group = train.groupby(['date_block_num', 'item_id']).agg({'item_price':
    ↪['mean']})
group.columns = ['avg_item_price_month']
group.reset_index(inplace=True)

#add to df
```

```
sales = pd.merge(sales, group, on=['date_block_num', 'item_id'], how='left')
sales['avg_item_price_month'] = sales['avg_item_price_month'].astype(np.float16)
```

```
[68]: #lag for price trend
lags = [1,2,3,4,5,6]
sales = lag_feature(sales, lags, 'avg_item_price_month')
```

```
[69]: for i in lags:
        sales['delta_price_lag_'+str(i)] = \
            (sales['avg_item_price_month_lag_'+str(i)] - sales['avg_item_price']) / \
            sales['avg_item_price']
```

```
[70]: def select_trend(row):
        for i in lags:
            if row['delta_price_lag_'+str(i)]:
                return row['delta_price_lag_'+str(i)]
        return 0
```

```
[71]: sales['delta_price_lag'] = sales.apply(select_trend, axis=1)
sales['delta_price_lag'] = sales['delta_price_lag'].astype(np.float16)
sales['delta_price_lag'].fillna(0, inplace=True)

features_to_drop = ['avg_item_price_month', 'avg_item_price']
for i in lags:
    features_to_drop += ['avg_item_price_month_lag_'+str(i)]
    features_to_drop += ['delta_price_lag_'+str(i)]
```

```
[72]: sales.drop(features_to_drop, axis=1, inplace=True)
```

```
[73]: sales.to_pickle('data_2.pickle.gzde', compression='gzip')
```

```
[74]: del sales
```

1.2.11 3. Shop revenue trends

```
[75]: sales = pd.read_pickle('data_2.pickle.gzde', compression='gzip')
```

```
[76]: #total revenue per shop per month
train['revenue'] = train['item_price'] * train['item_cnt_day']
group = train.groupby(['date_block_num', 'shop_id']).agg({'revenue': ['sum']})
group.columns = ['shop_revenue_month']
group.reset_index(inplace=True)

sales = pd.merge(sales, group, on=['date_block_num', 'shop_id'], how='left')
sales['shop_revenue_month'] = sales['shop_revenue_month'].astype(np.float32)
```

```
[77]: #average revenue per shop per month
group = group.groupby(['shop_id']).agg({'shop_revenue_month': ['mean']})
group.columns = ['shop_avg_revenue']
group.reset_index(inplace=True)

sales = pd.merge(sales, group, on=['shop_id'], how='left')
sales['shop_avg_revenue'] = sales['shop_avg_revenue'].astype(np.float32)

sales['delta_revenue'] = (sales['shop_revenue_month'] -
    ↪ sales['shop_avg_revenue']) / sales['shop_avg_revenue']

sales['delta_revenue'] = sales['delta_revenue'].astype(np.float16)

[78]: #lag revenue features
sales = lag_feature(sales, [1], 'delta_revenue')

[79]: #drop present month revenue features
sales.drop(['shop_revenue_month', 'shop_avg_revenue', 'delta_revenue'], axis=1,
    ↪ inplace=True)

[80]: sales.to_pickle('data_3.pickle.gzde', compression='gzip')

[81]: del sales
```

1.2.12 4. Date features

```
[82]: sales = pd.read_pickle('data_3.pickle.gzde', compression='gzip')

[83]: dates_train = train[['date', 'date_block_num']].drop_duplicates()
dates_test = dates_train[dates_train['date_block_num'] == 34-12]
dates_test['date_block_num'] = 34
dates_test['date'] = dates_test['date'] + pd.DateOffset(years=1)
dates_all = pd.concat([dates_train, dates_test])

dates_all['dow'] = dates_all['date'].dt.dayofweek
dates_all['year'] = dates_all['date'].dt.year
dates_all['month'] = dates_all['date'].dt.month
dates_all = pd.get_dummies(dates_all, columns=['dow'])
dow_col = ['dow_' + str(x) for x in range(7)]
date_features = dates_all.groupby(['year', 'month', 'date_block_num'])[dow_col].
    ↪ agg('sum').reset_index()
date_features['days_of_month'] = date_features[dow_col].sum(axis=1)
date_features['year'] = date_features['year'] - 2013

date_features = date_features[['month', 'year', 'days_of_month',
    ↪ 'date_block_num']]
```



```
sales = sales.merge(date_features, on = 'date_block_num', how = 'left')
date_columns = date_features.columns.difference(set(index_cols))
```

```
[84]: sales.head()
```

```
[84]:
```

	cat_subtype_id	cat_type_id	city_id	date_block_num	item_category_id	\
0	4	11	0	0	40	
1	10	5	0	0	19	
2	55	8	0	0	30	
3	16	5	0	0	23	
4	4	11	0	0	40	

	item_cnt_month	item_id	item_subtype_id	item_type_id	shop_id	\
0	0.0	19	42	4	2	
1	1.0	27	42	77	2	
2	0.0	28	42	108	2	
3	0.0	29	42	124	2	
4	0.0	32	42	4	2	

	shop_item_id	shop_type_id	item_cnt_month_lag_1	item_cnt_month_lag_2	\
0	2_19	5	NaN	NaN	
1	2_27	5	NaN	NaN	
2	2_28	5	NaN	NaN	
3	2_29	5	NaN	NaN	
4	2_32	5	NaN	NaN	

	item_cnt_month_lag_3	item_cnt_month_lag_4	item_cnt_month_lag_5	\
0	NaN	NaN	NaN	
1	NaN	NaN	NaN	
2	NaN	NaN	NaN	
3	NaN	NaN	NaN	
4	NaN	NaN	NaN	

	item_cnt_month_lag_6	item_cnt_month_lag_7	item_cnt_month_lag_8	\
0	NaN	NaN	NaN	
1	NaN	NaN	NaN	
2	NaN	NaN	NaN	
3	NaN	NaN	NaN	
4	NaN	NaN	NaN	

	item_cnt_month_lag_9	item_cnt_month_lag_10	item_cnt_month_lag_11	\
0	NaN	NaN	NaN	
1	NaN	NaN	NaN	
2	NaN	NaN	NaN	
3	NaN	NaN	NaN	
4	NaN	NaN	NaN	

	item_cnt_month_lag_12	shop_id_cnt_month_mean_Expanding_lag_1	...	\
0	NaN		NaN	...
1	NaN		NaN	...
2	NaN		NaN	...
3	NaN		NaN	...
4	NaN		NaN	...

	item_type_id_cnt_month_mean_Expanding_lag_1	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	
4	NaN	

	item_type_id_cnt_month_mean_Expanding_lag_2	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	
4	NaN	

	item_type_id_cnt_month_mean_Expanding_lag_3	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	
4	NaN	

	item_type_id_cnt_month_mean_Expanding_lag_4	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	
4	NaN	

	item_type_id_cnt_month_mean_Expanding_lag_12	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	
4	NaN	

	item_subtype_id_cnt_month_mean_Expanding_lag_1	\
0	NaN	
1	NaN	
2	NaN	
3	NaN	

4	NaN
---	-----

	item_subtype_id_cnt_month_mean_Expanding_lag_2 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_subtype_id_cnt_month_mean_Expanding_lag_3 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_subtype_id_cnt_month_mean_Expanding_lag_4 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	item_subtype_id_cnt_month_mean_Expanding_lag_12 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_type_id_cnt_month_mean_Expanding_lag_1 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_type_id_cnt_month_mean_Expanding_lag_2 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_type_id_cnt_month_mean_Expanding_lag_3 \
0	NaN
1	NaN

2	NaN
3	NaN
4	NaN

	cat_type_id_cnt_month_mean_Expanding_lag_4 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_type_id_cnt_month_mean_Expanding_lag_12 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_subtype_id_cnt_month_mean_Expanding_lag_1 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_subtype_id_cnt_month_mean_Expanding_lag_2 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_subtype_id_cnt_month_mean_Expanding_lag_3 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_subtype_id_cnt_month_mean_Expanding_lag_4 \
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

	cat_subtype_id_cnt_month_mean_Expanding_lag_12	delta_price_lag \
--	--	-------------------

0	NaN	0.0
1	NaN	0.0
2	NaN	0.0
3	NaN	0.0
4	NaN	0.0

	delta_revenue_lag_1	month	year	days_of_month
0	NaN	1	0	31
1	NaN	1	0	31
2	NaN	1	0	31
3	NaN	1	0	31
4	NaN	1	0	31

[5 rows x 79 columns]

1.2.13 5. Months since the last sale and first sale for each shop-item combination

Each row will be iterated through, treating $\{\text{shop_id}, \text{item_id}\}$ as the key and date_block_num as values. If the key is not in the cache and not equal to 0, the key-value pair is added to the cache. If the key is in the cache already, the difference between the current and previous date_block_num value is calculated.

These will need to be lagged as if the present month is the first time sales for a shop-item combination are recorded, it will be reported as 0 months, which would not be possible for sales that have yet to be recorded for prediction purposes.

```
[85]: #months since last sale of shop-item combination
cache = {}
sales['months_since_item_shop_last_sale'] = -1
sales['months_since_item_shop_last_sale'] = 0
sales['months_since_item_shop_last_sale'] = sales['months_since_item_shop_last_sale'].astype(np.int8)
for idx, row in sales.iterrows():
    key = str(row.item_id)+' '+str(row.shop_id)
    if key not in cache:
        if row.item_cnt_month!=0:
            cache[key] = row.date_block_num
    else:
        last_date_block_num = cache[key]
        sales.at[idx, 'months_since_item_shop_last_sale'] = row.date_block_num - last_date_block_num
        cache[key] = row.date_block_num
```

```
[86]: train = pd.read_pickle('train_alt.pickle.gz', compression='gzip')
```

```
[87]: group = train.groupby(['date_block_num', 'shop_id', 'item_id']).
        .agg({'item_cnt_day': ['sum']})
        group.columns = ['item_cnt_month']
```

```
group.reset_index(inplace=True)
```

```
[88]: #months since first sale of shop-item combination
group['months_since_shop_item_first_sale'] = group['date_block_num'] - group.
    ↳groupby(['shop_id', 'item_id'])['date_block_num'].transform('min')
group1 =_
    ↳group[['date_block_num', 'shop_id', 'item_id', 'months_since_shop_item_first_sale']]
```

```
[89]: sales = pd.merge(sales, group1, on=['date_block_num', 'shop_id', 'item_id'],_
    ↳how='left')
sales['months_since_shop_item_first_sale'] =_
    ↳sales['months_since_shop_item_first_sale'].fillna(-1)
```

```
[90]: #lag features
sales = lag_feature(sales, [1], 'months_since_item_shop_last_sale')
sales = lag_feature(sales, [1], 'months_since_shop_item_first_sale')

sales = sales.
    ↳drop(['months_since_item_shop_last_sale', 'months_since_shop_item_first_sale'],_
    ↳axis=1)
```

```
[91]: sales.to_pickle('data_4.pickle.gzde', compression='gzip')
```

```
[92]: del sales
```

1.2.14 6. Months since the last and first sale for each item only

The same approach is used as above, but with the key only being item_id.

```
[93]: sales = pd.read_pickle('data_4.pickle.gzde', compression='gzip')
```

```
[94]: sales = downcast_dtypes(sales)
```

```
[95]: #months since last sale of item
cache = {}
sales['months_since_item_last_sale'] = -1
sales['months_since_item_last_sale'] = sales['months_since_item_last_sale'].
    ↳astype(np.int8)
for idx, row in sales.iterrows():
    key = row.item_id
    if key not in cache:
        if row.item_cnt_month!=0:
            cache[key] = row.date_block_num
    else:
        last_date_block_num = cache[key]
        if row.date_block_num>last_date_block_num:
```

```

        sales.at[idx, 'months_since_item_last_sale'] = row.date_block_num -
        ↪last_date_block_num
        cache[key] = row.date_block_num

sales = lag_feature(sales, [1], 'months_since_item_last_sale')
sales = sales.drop(['months_since_item_last_sale'], axis=1)
sales = downcast_dtypes(sales)

```

```
[96]: train = pd.read_pickle('train_alt.pickle.gzde', compression='gzip')
```

```
[97]: group = train.groupby(['date_block_num', 'shop_id', 'item_id']).
        ↪agg({'item_cnt_day': ['sum']})
group.columns = ['item_cnt_month']
group.reset_index(inplace=True)

#months since first sale of item
group['months_since_item_first_sale'] = group['date_block_num'] - group.
        ↪groupby(['item_id'])['date_block_num'].transform('min')
group =
        ↪group[['date_block_num', 'shop_id', 'item_id', 'months_since_item_first_sale']]

sales = pd.merge(sales, group, on=['date_block_num', 'shop_id', 'item_id'],
        ↪how='left')
sales['months_since_item_first_sale'] = sales['months_since_item_first_sale'].
        ↪fillna(-1)
sales = lag_feature(sales, [1], 'months_since_item_first_sale')
sales = sales.drop(['months_since_item_first_sale'], axis=1)

```

1.2.15 Drop first 12 months of data as lags incomplete

```
[98]: sales = sales[sales.date_block_num > 11]
```

1.2.16 Fill remaining null values with zero

```
[99]: sales = sales.fillna(0)
sales = downcast_dtypes(sales)
sales.head()
```

```
[99]:
```

	cat_subtype_id	cat_type_id	city_id	date_block_num	\
3395293	4	11	0	12	
3395294	1	11	0	12	
3395295	4	11	0	12	
3395296	1	11	0	12	
3395297	4	11	0	12	

	item_category_id	item_cnt_month	item_id	item_subtype_id	\
3395293	40	0.0	30	42	
3395294	37	0.0	31	562	
3395295	40	1.0	32	42	
3395296	37	1.0	33	562	
3395297	40	0.0	34	1367	

	item_type_id	shop_id	shop_item_id	shop_type_id	\
3395293	4	2	2_30	5	
3395294	4	2	2_31	5	
3395295	4	2	2_32	5	
3395296	4	2	2_33	5	
3395297	4	2	2_34	5	

	item_cnt_month_lag_1	item_cnt_month_lag_2	item_cnt_month_lag_3	\
3395293	0.0	0.0	0.0	
3395294	0.0	0.0	0.0	
3395295	0.0	0.0	0.0	
3395296	1.0	2.0	0.0	
3395297	0.0	0.0	0.0	

	item_cnt_month_lag_4	item_cnt_month_lag_5	item_cnt_month_lag_6	\
3395293	0.0	0.0	0.0	
3395294	0.0	0.0	0.0	
3395295	0.0	0.0	0.0	
3395296	0.0	0.0	0.0	
3395297	0.0	0.0	0.0	

	item_cnt_month_lag_7	item_cnt_month_lag_8	item_cnt_month_lag_9	\
3395293	1.0	0.0	0.0	
3395294	0.0	0.0	1.0	
3395295	0.0	0.0	0.0	
3395296	0.0	0.0	0.0	
3395297	0.0	0.0	0.0	

	item_cnt_month_lag_10	item_cnt_month_lag_11	item_cnt_month_lag_12	\
3395293	1.0	0.0	0.0	
3395294	1.0	4.0	0.0	
3395295	0.0	0.0	0.0	
3395296	0.0	0.0	1.0	
3395297	0.0	0.0	0.0	

	shop_id_cnt_month_mean_Expanding_lag_1	...	\
3395293	0.099054	...	
3395294	0.099053	...	
3395295	0.099052	...	
3395296	0.099051	...	

3395297	0.099062	...
item_type_id_cnt_month_mean_Expanding_lag_12 \		
3395293	0.000000	
3395294	0.000000	
3395295	0.000000	
3395296	0.000000	
3395297	0.333333	
item_subtype_id_cnt_month_mean_Expanding_lag_1 \		
3395293	0.384450	
3395294	0.201874	
3395295	0.384450	
3395296	0.201874	
3395297	0.192280	
item_subtype_id_cnt_month_mean_Expanding_lag_2 \		
3395293	0.382351	
3395294	0.204375	
3395295	0.382351	
3395296	0.204375	
3395297	0.194044	
item_subtype_id_cnt_month_mean_Expanding_lag_3 \		
3395293	0.382264	
3395294	0.207749	
3395295	0.382264	
3395296	0.207748	
3395297	0.195743	
item_subtype_id_cnt_month_mean_Expanding_lag_4 \		
3395293	0.379615	
3395294	0.207921	
3395295	0.379615	
3395296	0.207920	
3395297	0.199872	
item_subtype_id_cnt_month_mean_Expanding_lag_12 \		
3395293	0.000000	
3395294	0.000000	
3395295	0.250000	
3395296	0.311493	
3395297	0.311493	
cat_type_id_cnt_month_mean_Expanding_lag_1 \		
3395293	0.238131	
3395294	0.238130	

3395295	0.238130
3395296	0.238130
3395297	0.238131

	cat_type_id_cnt_month_mean_Expanding_lag_2 \
3395293	0.238516
3395294	0.238516
3395295	0.238516
3395296	0.238515
3395297	0.238517

	cat_type_id_cnt_month_mean_Expanding_lag_3 \
3395293	0.240038
3395294	0.240038
3395295	0.240038
3395296	0.240037
3395297	0.240037

	cat_type_id_cnt_month_mean_Expanding_lag_4 \
3395293	0.242356
3395294	0.242356
3395295	0.242356
3395296	0.242355
3395297	0.242355

	cat_type_id_cnt_month_mean_Expanding_lag_12 \
3395293	0.000000
3395294	0.000000
3395295	0.000000
3395296	0.000000
3395297	0.333333

	cat_subtype_id_cnt_month_mean_Expanding_lag_1 \
3395293	0.254690
3395294	0.205493
3395295	0.254689
3395296	0.205492
3395297	0.254689

	cat_subtype_id_cnt_month_mean_Expanding_lag_2 \
3395293	0.254084
3395294	0.207230
3395295	0.254084
3395296	0.207229
3395297	0.254084

	cat_subtype_id_cnt_month_mean_Expanding_lag_3 \
--	---

3395293	0.254229
3395294	0.210747
3395295	0.254229
3395296	0.210746
3395297	0.254229

cat_subtype_id_cnt_month_mean_Expanding_lag_4 \	
3395293	0.256372
3395294	0.212243
3395295	0.256372
3395296	0.212243
3395297	0.256371

cat_subtype_id_cnt_month_mean_Expanding_lag_12 delta_price_lag \		
3395293	0.000000	-0.478760
3395294	0.000000	-0.127563
3395295	0.000000	-0.407959
3395296	0.311493	-0.218750
3395297	0.000000	0.005058

delta_revenue_lag_1 month year days_of_month \				
3395293	1.211914	1	1	31
3395294	1.211914	1	1	31
3395295	1.211914	1	1	31
3395296	1.211914	1	1	31
3395297	1.211914	1	1	31

months_since_item_shop_last_sale_lag_1 \	
3395293	1.0
3395294	1.0
3395295	-1.0
3395296	1.0
3395297	-1.0

months_since_shop_item_first_sale_lag_1 \	
3395293	-1.0
3395294	-1.0
3395295	-1.0
3395296	11.0
3395297	-1.0

months_since_item_last_sale_lag_1 months_since_item_first_sale_lag_1		
3395293	1.0	-1.0
3395294	1.0	-1.0
3395295	1.0	-1.0
3395296	1.0	11.0
3395297	1.0	-1.0

[5 rows x 83 columns]

```
[100]: sales.describe().round(2)
```

```
[100]:
```

	cat_subtype_id	cat_type_id	city_id	date_block_num	\
count	5465978.00	5465978.00	5465978.00	5465978.00	
mean	19.74	10.96	15.38	22.53	
std	20.96	3.01	8.29	6.60	
min	0.00	0.00	0.00	12.00	
25%	2.00	11.00	10.00	17.00	
50%	10.00	11.00	15.00	22.00	
75%	35.00	13.00	22.00	28.00	
max	64.00	19.00	30.00	34.00	

	item_category_id	item_cnt_month	item_id	item_subtype_id	\
count	5465978.00	5465978.00	5465978.00	5465978.00	
mean	44.86	0.29	11211.40	404.73	
std	15.84	1.16	6268.25	503.61	
min	0.00	0.00	1.00	0.00	
25%	37.00	0.00	5647.00	42.00	
50%	40.00	0.00	11319.00	42.00	
75%	55.00	0.00	16461.00	637.00	
max	83.00	20.00	22169.00	1665.00	

	item_type_id	shop_id	shop_type_id	item_cnt_month_lag_1	\
count	5465978.00	5465978.00	5465978.00	5465978.00	
mean	20.14	31.45	3.58	0.30	
std	35.02	17.83	1.80	1.21	
min	0.00	2.00	0.00	0.00	
25%	4.00	16.00	3.00	0.00	
50%	4.00	34.00	4.00	0.00	
75%	4.00	47.00	5.00	0.00	
max	174.00	59.00	5.00	20.00	

	item_cnt_month_lag_2	item_cnt_month_lag_3	item_cnt_month_lag_4	\
count	5465978.00	5465978.00	5465978.00	
mean	0.30	0.30	0.30	
std	1.22	1.22	1.24	
min	0.00	0.00	0.00	
25%	0.00	0.00	0.00	
50%	0.00	0.00	0.00	
75%	0.00	0.00	0.00	
max	20.00	20.00	20.00	

	item_cnt_month_lag_5	item_cnt_month_lag_6	item_cnt_month_lag_7	\
count	5465978.00	5465978.00	5465978.00	

mean	0.30	0.29	0.29
std	1.24	1.24	1.23
min	0.00	0.00	0.00
25%	0.00	0.00	0.00
50%	0.00	0.00	0.00
75%	0.00	0.00	0.00
max	20.00	20.00	20.00

	item_cnt_month_lag_8	item_cnt_month_lag_9	item_cnt_month_lag_10 \
count	5465978.00	5465978.00	5465978.00
mean	0.28	0.27	0.27
std	1.22	1.21	1.21
min	0.00	0.00	0.00
25%	0.00	0.00	0.00
50%	0.00	0.00	0.00
75%	0.00	0.00	0.00
max	20.00	20.00	20.00

	item_cnt_month_lag_11	item_cnt_month_lag_12 \
count	5465978.00	5465978.00
mean	0.26	0.25
std	1.20	1.17
min	0.00	0.00
25%	0.00	0.00
50%	0.00	0.00
75%	0.00	0.00
max	20.00	20.00

	shop_id_cnt_month_mean_Expanding_lag_1 \
count	5465978.00
mean	0.25
std	0.23
min	0.00
25%	0.11
50%	0.22
75%	0.29
max	1.21

	shop_id_cnt_month_mean_Expanding_lag_2 ... \
count	5465978.00 ...
mean	0.24 ...
std	0.23 ...
min	0.00 ...
25%	0.04 ...
50%	0.21 ...
75%	0.29 ...
max	1.21 ...

	item_type_id_cnt_month_mean_Expanding_lag_12 \
count	5465978.00
mean	0.17
std	0.34
min	0.00
25%	0.00
50%	0.00
75%	0.24
max	19.84

	item_subtype_id_cnt_month_mean_Expanding_lag_1 \
count	5465978.00
mean	0.28
std	0.42
min	0.00
25%	0.10
50%	0.28
75%	0.39
max	20.00

	item_subtype_id_cnt_month_mean_Expanding_lag_2 \
count	5465978.00
mean	0.27
std	0.43
min	0.00
25%	0.04
50%	0.21
75%	0.39
max	20.00

	item_subtype_id_cnt_month_mean_Expanding_lag_3 \
count	5465978.00
mean	0.26
std	0.43
min	0.00
25%	0.00
50%	0.20
75%	0.39
max	20.00

	item_subtype_id_cnt_month_mean_Expanding_lag_4 \
count	5465978.00
mean	0.25
std	0.43
min	0.00
25%	0.00

50%	0.20
75%	0.39
max	20.00

	item_subtype_id_cnt_month_mean_Expanding_lag_12 \
count	5465978.00
mean	0.17
std	0.42
min	0.00
25%	0.00
50%	0.00
75%	0.38
max	20.00

	cat_type_id_cnt_month_mean_Expanding_lag_1 \
count	5465978.00
mean	0.27
std	0.27
min	0.00
25%	0.13
50%	0.23
75%	0.30
max	6.39

	cat_type_id_cnt_month_mean_Expanding_lag_2 \
count	5465978.00
mean	0.26
std	0.27
min	0.00
25%	0.12
50%	0.23
75%	0.30
max	6.39

	cat_type_id_cnt_month_mean_Expanding_lag_3 \
count	5465978.00
mean	0.25
std	0.27
min	0.00
25%	0.00
50%	0.23
75%	0.29
max	6.39

	cat_type_id_cnt_month_mean_Expanding_lag_4 \
count	5465978.00
mean	0.24

std	0.27
min	0.00
25%	0.00
50%	0.18
75%	0.28
max	6.39

cat_type_id_cnt_month_mean_Expanding_lag_12 \	
count	5465978.00
mean	0.16
std	0.26
min	0.00
25%	0.00
50%	0.00
75%	0.24
max	8.00

cat_subtype_id_cnt_month_mean_Expanding_lag_1 \	
count	5465978.00
mean	0.29
std	0.44
min	0.00
25%	0.07
50%	0.22
75%	0.26
max	15.33

cat_subtype_id_cnt_month_mean_Expanding_lag_2 \	
count	5465978.00
mean	0.28
std	0.43
min	0.00
25%	0.05
50%	0.21
75%	0.26
max	15.33

cat_subtype_id_cnt_month_mean_Expanding_lag_3 \	
count	5465978.00
mean	0.27
std	0.43
min	0.00
25%	0.00
50%	0.21
75%	0.26
max	15.33

	cat_subtype_id_cnt_month_mean_Expanding_lag_4 \
count	5465978.00
mean	0.25
std	0.42
min	0.00
25%	0.00
50%	0.20
75%	0.26
max	15.33

	cat_subtype_id_cnt_month_mean_Expanding_lag_12	delta_price_lag \
count	5465978.00	5465978.00
mean	0.17	NaN
std	0.36	0.00
min	0.00	-1.00
25%	0.00	-0.04
50%	0.00	0.00
75%	0.22	0.01
max	15.00	2.99

	delta_revenue_lag_1	month	year	days_of_month \
count	5465978.00	5465978.00	5465978.00	5465978.00
mean	NaN	6.18	1.45	30.39
std	0.00	3.37	0.50	0.88
min	-1.00	1.00	1.00	28.00
25%	-0.18	3.00	1.00	30.00
50%	0.00	6.00	1.00	31.00
75%	0.07	9.00	2.00	31.00
max	4.20	12.00	2.00	31.00

	months_since_item_shop_last_sale_lag_1 \
count	5465978.00
mean	0.17
std	1.01
min	-1.00
25%	-1.00
50%	0.00
75%	1.00
max	31.00

	months_since_shop_item_first_sale_lag_1 \
count	5465978.00
mean	0.65
std	4.58
min	-1.00
25%	-1.00
50%	-1.00

75%	0.00
max	33.00

	months_since_item_last_sale_lag_1	months_since_item_first_sale_lag_1
count	5465978.00	5465978.00
mean	-0.75	0.93
std	0.51	5.12
min	-1.00	-1.00
25%	-1.00	-1.00
50%	-1.00	-1.00
75%	-1.00	0.00
max	31.00	33.00

[8 rows x 82 columns]

```
[101]: sales = sales.drop(['shop_item_id'], axis=1)
```

```
[102]: sales.to_pickle('data_5.pickle.gzde', compression='gzip')
```

```
[103]: del sales
del group
del cache
```

1.2.17 Scaling data

The algorithms used to model are tree-based, therefore the data does not require normalisation.

```
[104]: sales = pd.read_pickle('data_5.pickle.gzde', compression='gzip')
```

```
[105]: from sklearn.preprocessing import MinMaxScaler
mms = MinMaxScaler()

feature_cols = list(sales)
feature_cols = [e for e in feature_cols if e not in ['date_block_num', 'item_cnt_month']]

train = sales[sales['date_block_num'] != sales['date_block_num'].max()]
test = sales[sales['date_block_num'] == sales['date_block_num'].max()]
```

```
[106]: train[feature_cols] = mms.fit_transform(train[feature_cols])
test[feature_cols] = mms.transform(test[feature_cols])
```

```
[107]: sales = pd.concat([train, test], axis = 0)
sales = downcast_dtypes(sales)
```

```
[108]: del train, test, feature_cols
gc.collect()
```

[108]: 20

1.2.18 Feature selection

Multiple methods will be used to find the most important features in deriving a model to predict sales. The most consistently important features will be selected for the purposes of building the model.

```
[109]: feat_sel = sales[sales.date_block_num < 33]
X = feat_sel.drop(['item_cnt_month', 'date_block_num'], axis=1)
Y = feat_sel['item_cnt_month']
feature_name = list(X.columns)
print(len(feature_name))
print(feature_name)
```

80

```
['cat_subtype_id', 'cat_type_id', 'city_id', 'item_category_id', 'item_id',
'item_subtype_id', 'item_type_id', 'shop_id', 'shop_type_id',
'item_cnt_month_lag_1', 'item_cnt_month_lag_2', 'item_cnt_month_lag_3',
'item_cnt_month_lag_4', 'item_cnt_month_lag_5', 'item_cnt_month_lag_6',
'item_cnt_month_lag_7', 'item_cnt_month_lag_8', 'item_cnt_month_lag_9',
'item_cnt_month_lag_10', 'item_cnt_month_lag_11', 'item_cnt_month_lag_12',
'shop_id_cnt_month_mean_Expanding_lag_1',
'shop_id_cnt_month_mean_Expanding_lag_2',
'shop_id_cnt_month_mean_Expanding_lag_3',
'shop_id_cnt_month_mean_Expanding_lag_4',
'shop_id_cnt_month_mean_Expanding_lag_12',
'item_id_cnt_month_mean_Expanding_lag_1',
'item_id_cnt_month_mean_Expanding_lag_2',
'item_id_cnt_month_mean_Expanding_lag_3',
'item_id_cnt_month_mean_Expanding_lag_4',
'item_id_cnt_month_mean_Expanding_lag_12',
'shop_item_id_cnt_month_mean_Smooth_lag_1',
'shop_item_id_cnt_month_mean_Smooth_lag_2',
'shop_item_id_cnt_month_mean_Smooth_lag_3',
'shop_item_id_cnt_month_mean_Smooth_lag_4',
'shop_item_id_cnt_month_mean_Smooth_lag_12',
'shop_type_id_cnt_month_mean_Expanding_lag_1',
'shop_type_id_cnt_month_mean_Expanding_lag_2',
'shop_type_id_cnt_month_mean_Expanding_lag_3',
'shop_type_id_cnt_month_mean_Expanding_lag_4',
'shop_type_id_cnt_month_mean_Expanding_lag_12',
'city_id_cnt_month_mean_Expanding_lag_1',
'city_id_cnt_month_mean_Expanding_lag_2',
'city_id_cnt_month_mean_Expanding_lag_3',
'city_id_cnt_month_mean_Expanding_lag_4',
'city_id_cnt_month_mean_Expanding_lag_12',
'item_category_id_cnt_month_mean_Expanding_lag_1',
```

```

'item_category_id_cnt_month_mean_Expanding_lag_2',
'item_category_id_cnt_month_mean_Expanding_lag_3',
'item_category_id_cnt_month_mean_Expanding_lag_4',
'item_category_id_cnt_month_mean_Expanding_lag_12',
'item_type_id_cnt_month_mean_Expanding_lag_1',
'item_type_id_cnt_month_mean_Expanding_lag_2',
'item_type_id_cnt_month_mean_Expanding_lag_3',
'item_type_id_cnt_month_mean_Expanding_lag_4',
'item_type_id_cnt_month_mean_Expanding_lag_12',
'item_subtype_id_cnt_month_mean_Expanding_lag_1',
'item_subtype_id_cnt_month_mean_Expanding_lag_2',
'item_subtype_id_cnt_month_mean_Expanding_lag_3',
'item_subtype_id_cnt_month_mean_Expanding_lag_4',
'item_subtype_id_cnt_month_mean_Expanding_lag_12',
'cat_type_id_cnt_month_mean_Expanding_lag_1',
'cat_type_id_cnt_month_mean_Expanding_lag_2',
'cat_type_id_cnt_month_mean_Expanding_lag_3',
'cat_type_id_cnt_month_mean_Expanding_lag_4',
'cat_type_id_cnt_month_mean_Expanding_lag_12',
'cat_subtype_id_cnt_month_mean_Expanding_lag_1',
'cat_subtype_id_cnt_month_mean_Expanding_lag_2',
'cat_subtype_id_cnt_month_mean_Expanding_lag_3',
'cat_subtype_id_cnt_month_mean_Expanding_lag_4',
'cat_subtype_id_cnt_month_mean_Expanding_lag_12', 'delta_price_lag',
'delta_revenue_lag_1', 'month', 'year', 'days_of_month',
'months_since_item_shop_last_sale_lag_1',
'months_since_shop_item_first_sale_lag_1', 'months_since_item_last_sale_lag_1',
'months_since_item_first_sale_lag_1']

```

```
[110]: num_feats = 30
```

```

[111]: #pearson correlation
def cor_selector(X, y,num_feats):
    cor_list = []
    feature_name = X.columns.tolist()
    # calculate the correlation with y for each feature
    for i in X.columns.tolist():
        cor = np.corrcoef(X[i], y)[0, 1]
        cor_list.append(cor)
    # replace NaN with 0
    cor_list = [0 if np.isnan(i) else i for i in cor_list]
    # feature name
    cor_feature = X.iloc[:,np.argsort(np.abs(cor_list))[-num_feats:]].columns.
    ↪tolist()
    # feature selection? 0 for not select, 1 for select
    cor_support = [True if i in cor_feature else False for i in feature_name]
    return cor_support, cor_feature

```

```
cor_support, cor_feature = cor_selector(X, Y, num_feats)
print(str(len(cor_feature)), 'selected features')
```

30 selected features

```
[112]: cor_feature
```

```
[112]: ['item_subtype_id_cnt_month_mean_Expanding_lag_2',
        'shop_item_id_cnt_month_mean_Smooth_lag_12',
        'item_cnt_month_lag_12',
        'cat_subtype_id_cnt_month_mean_Expanding_lag_3',
        'item_category_id_cnt_month_mean_Expanding_lag_3',
        'months_since_shop_item_first_sale_lag_1',
        'item_cnt_month_lag_11',
        'item_subtype_id_cnt_month_mean_Expanding_lag_1',
        'item_cnt_month_lag_10',
        'cat_subtype_id_cnt_month_mean_Expanding_lag_2',
        'item_category_id_cnt_month_mean_Expanding_lag_2',
        'item_cnt_month_lag_9',
        'item_id_cnt_month_mean_Expanding_lag_4',
        'item_cnt_month_lag_8',
        'cat_subtype_id_cnt_month_mean_Expanding_lag_1',
        'item_category_id_cnt_month_mean_Expanding_lag_1',
        'item_cnt_month_lag_7',
        'item_id_cnt_month_mean_Expanding_lag_3',
        'item_cnt_month_lag_6',
        'item_cnt_month_lag_5',
        'shop_item_id_cnt_month_mean_Smooth_lag_4',
        'item_id_cnt_month_mean_Expanding_lag_2',
        'item_cnt_month_lag_4',
        'shop_item_id_cnt_month_mean_Smooth_lag_3',
        'shop_item_id_cnt_month_mean_Smooth_lag_2',
        'item_cnt_month_lag_3',
        'item_id_cnt_month_mean_Expanding_lag_1',
        'shop_item_id_cnt_month_mean_Smooth_lag_1',
        'item_cnt_month_lag_2',
        'item_cnt_month_lag_1']
```

```
[113]: #chi-squared
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
chi_selector = SelectKBest(chi2, k=num_feats)
chi_selector.fit(X, Y)
chi_support = chi_selector.get_support()
chi_feature = X.loc[:,chi_support].columns.tolist()
print(str(len(chi_feature)), 'selected features')
```

30 selected features

```
[114]: chi_feature
```

```
[114]: ['item_cnt_month_lag_1',
        'item_cnt_month_lag_2',
        'item_cnt_month_lag_3',
        'item_cnt_month_lag_4',
        'item_cnt_month_lag_5',
        'item_cnt_month_lag_6',
        'item_cnt_month_lag_7',
        'item_cnt_month_lag_8',
        'item_cnt_month_lag_9',
        'item_cnt_month_lag_10',
        'item_cnt_month_lag_11',
        'item_cnt_month_lag_12',
        'shop_id_cnt_month_mean_Expanding_lag_1',
        'shop_id_cnt_month_mean_Expanding_lag_2',
        'shop_id_cnt_month_mean_Expanding_lag_3',
        'shop_id_cnt_month_mean_Expanding_lag_4',
        'item_id_cnt_month_mean_Expanding_lag_1',
        'item_id_cnt_month_mean_Expanding_lag_2',
        'item_id_cnt_month_mean_Expanding_lag_3',
        'item_id_cnt_month_mean_Expanding_lag_4',
        'item_id_cnt_month_mean_Expanding_lag_12',
        'shop_item_id_cnt_month_mean_Smooth_lag_1',
        'shop_item_id_cnt_month_mean_Smooth_lag_2',
        'shop_item_id_cnt_month_mean_Smooth_lag_3',
        'shop_item_id_cnt_month_mean_Smooth_lag_4',
        'city_id_cnt_month_mean_Expanding_lag_1',
        'city_id_cnt_month_mean_Expanding_lag_2',
        'city_id_cnt_month_mean_Expanding_lag_3',
        'months_since_shop_item_first_sale_lag_1',
        'months_since_item_first_sale_lag_1']
```

```
[115]: #recursive feature elimination with linear regression
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression

rfe_selector = RFE(estimator=LinearRegression(),
    ↪n_features_to_select=num_feats, step=10, verbose=5)
rfe_selector.fit(X, Y)
rfe_support = rfe_selector.get_support()
rfe_feature = X.loc[:,rfe_support].columns.tolist()
print(str(len(rfe_feature)), 'selected features')
```

Fitting estimator with 80 features.

Fitting estimator with 70 features.

Fitting estimator with 60 features.

Fitting estimator with 50 features.
Fitting estimator with 40 features.
30 selected features

```
[116]: rfe_feature
```

```
[116]: ['item_cnt_month_lag_1',  
        'item_cnt_month_lag_2',  
        'item_cnt_month_lag_3',  
        'item_id_cnt_month_mean_Expanding_lag_1',  
        'item_id_cnt_month_mean_Expanding_lag_2',  
        'shop_item_id_cnt_month_mean_Smooth_lag_1',  
        'shop_item_id_cnt_month_mean_Smooth_lag_2',  
        'shop_item_id_cnt_month_mean_Smooth_lag_3',  
        'shop_item_id_cnt_month_mean_Smooth_lag_4',  
        'shop_item_id_cnt_month_mean_Smooth_lag_12',  
        'shop_type_id_cnt_month_mean_Expanding_lag_1',  
        'shop_type_id_cnt_month_mean_Expanding_lag_12',  
        'item_category_id_cnt_month_mean_Expanding_lag_1',  
        'item_category_id_cnt_month_mean_Expanding_lag_2',  
        'item_category_id_cnt_month_mean_Expanding_lag_3',  
        'item_category_id_cnt_month_mean_Expanding_lag_4',  
        'item_category_id_cnt_month_mean_Expanding_lag_12',  
        'item_type_id_cnt_month_mean_Expanding_lag_1',  
        'item_type_id_cnt_month_mean_Expanding_lag_2',  
        'item_subtype_id_cnt_month_mean_Expanding_lag_1',  
        'item_subtype_id_cnt_month_mean_Expanding_lag_2',  
        'item_subtype_id_cnt_month_mean_Expanding_lag_3',  
        'item_subtype_id_cnt_month_mean_Expanding_lag_4',  
        'cat_type_id_cnt_month_mean_Expanding_lag_1',  
        'cat_subtype_id_cnt_month_mean_Expanding_lag_1',  
        'cat_subtype_id_cnt_month_mean_Expanding_lag_2',  
        'cat_subtype_id_cnt_month_mean_Expanding_lag_3',  
        'cat_subtype_id_cnt_month_mean_Expanding_lag_4',  
        'cat_subtype_id_cnt_month_mean_Expanding_lag_12',  
        'months_since_item_last_sale_lag_1']
```

```
[117]: #recursive feature elimination with random forest  
#the SelectFromModel method was attempted, but running time was too long  
from sklearn.ensemble import RandomForestRegressor  
  
rf = RandomForestRegressor(n_estimators = 100,  
                           n_jobs = -1,  
                           oob_score = True,  
                           bootstrap = True,  
                           random_state = 42)  
  
rf.fit(X, Y)
```

```

r_lst = []
for feature in sorted(zip(rf.feature_importances_, X.columns), reverse=True):
    r_lst.append(feature)
r_lst = r_lst[:30]
features_rf = []
for i in r_lst:
    features_rf.append(i[1])

```

```

[118]: # Feature importance with light GBM
#lightGBM model fit
gbm = lgb.LGBMRegressor()
gbm.fit(X, Y)
gbm.booster_.feature_importance()

g_lst = []
for feature in sorted(zip(gbm.feature_importances_, X.columns), reverse=True):
    g_lst.append(feature)
g_lst = g_lst[:30]
features_gbm = []
for i in g_lst:
    features_gbm.append(i[1])

```

```

[119]: xgbm = xgb.XGBRegressor()
xgbm.fit(X, Y)

x_lst = []
for feature in sorted(zip(xgbm.feature_importances_, X.columns), reverse=True):
    x_lst.append(feature)
x_lst

```

[21:48:58] WARNING: src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

[21:48:58] WARNING: src/learner.cc:686: Tree method is automatically selected to be 'approx' for faster speed. To use old behavior (exact greedy algorithm on single machine), set tree_method to 'exact'.

```

[119]: [(0.29189757, 'item_cnt_month_lag_1'),
(0.0972129, 'shop_item_id_cnt_month_mean_Smooth_lag_1'),
(0.07092963, 'cat_subtype_id_cnt_month_mean_Expanding_lag_4'),
(0.05207234, 'months_since_item_shop_last_sale_lag_1'),
(0.03970951, 'shop_type_id_cnt_month_mean_Expanding_lag_4'),
(0.032093115, 'item_cnt_month_lag_2'),
(0.030606989, 'months_since_item_first_sale_lag_1'),
(0.024894359, 'shop_type_id_cnt_month_mean_Expanding_lag_3'),
(0.02302277, 'shop_id_cnt_month_mean_Expanding_lag_12'),
(0.019293122, 'item_cnt_month_lag_3'),

```



```

(0.01909757, 'item_category_id'),
(0.018602934, 'item_cnt_month_lag_8'),
(0.017617425, 'months_since_shop_item_first_sale_lag_1'),
(0.01657576, 'item_type_id_cnt_month_mean_Expanding_lag_2'),
(0.015775928, 'shop_item_id_cnt_month_mean_Smooth_lag_2'),
(0.015747877, 'shop_item_id_cnt_month_mean_Smooth_lag_3'),
(0.015213964, 'shop_item_id_cnt_month_mean_Smooth_lag_4'),
(0.012944539, 'month'),
(0.011845271, 'item_subtype_id_cnt_month_mean_Expanding_lag_4'),
(0.011731102, 'shop_id_cnt_month_mean_Expanding_lag_4'),
(0.009351417, 'item_cnt_month_lag_5'),
(0.008769338, 'shop_item_id_cnt_month_mean_Smooth_lag_12'),
(0.007908101, 'cat_type_id_cnt_month_mean_Expanding_lag_3'),
(0.007446583, 'cat_type_id_cnt_month_mean_Expanding_lag_4'),
(0.007249106, 'cat_type_id_cnt_month_mean_Expanding_lag_12'),
(0.006701999, 'item_cnt_month_lag_6'),
(0.0065668696, 'item_id_cnt_month_mean_Expanding_lag_2'),
(0.005944074, 'item_id_cnt_month_mean_Expanding_lag_4'),
(0.0058902176, 'item_id_cnt_month_mean_Expanding_lag_12'),
(0.0057972954, 'city_id'),
(0.0056893537, 'item_type_id_cnt_month_mean_Expanding_lag_12'),
(0.00542173, 'shop_type_id_cnt_month_mean_Expanding_lag_2'),
(0.0049782908, 'cat_subtype_id_cnt_month_mean_Expanding_lag_3'),
(0.00486983, 'item_cnt_month_lag_10'),
(0.0046798512, 'item_type_id_cnt_month_mean_Expanding_lag_3'),
(0.0046731657, 'item_type_id'),
(0.0046636546, 'cat_subtype_id'),
(0.0046461974, 'cat_type_id_cnt_month_mean_Expanding_lag_1'),
(0.004508725, 'item_cnt_month_lag_7'),
(0.004317061, 'cat_type_id_cnt_month_mean_Expanding_lag_2'),
(0.0042478885, 'item_cnt_month_lag_4'),
(0.004143697, 'item_category_id_cnt_month_mean_Expanding_lag_1'),
(0.003965197, 'cat_type_id'),
(0.0037120741, 'item_id_cnt_month_mean_Expanding_lag_1'),
(0.0031850059, 'delta_price_lag'),
(0.0028749283, 'item_type_id_cnt_month_mean_Expanding_lag_1'),
(0.0028671175, 'item_subtype_id_cnt_month_mean_Expanding_lag_3'),
(0.0025235137, 'cat_subtype_id_cnt_month_mean_Expanding_lag_12'),
(0.0020568229, 'shop_id_cnt_month_mean_Expanding_lag_1'),
(0.0020163655, 'item_cnt_month_lag_9'),
(0.0020153865, 'item_cnt_month_lag_11'),
(0.0019541846, 'year'),
(0.0019099995, 'item_id_cnt_month_mean_Expanding_lag_3'),
(0.0015732666, 'days_of_month'),
(0.0014899498, 'item_subtype_id'),
(0.0013850734, 'item_cnt_month_lag_12'),
(0.0011220238, 'item_id'),

```

```
(0.0, 'shop_type_id_cnt_month_mean_Expanding_lag_12'),
(0.0, 'shop_type_id_cnt_month_mean_Expanding_lag_1'),
(0.0, 'shop_type_id'),
(0.0, 'shop_id_cnt_month_mean_Expanding_lag_3'),
(0.0, 'shop_id_cnt_month_mean_Expanding_lag_2'),
(0.0, 'shop_id'),
(0.0, 'months_since_item_last_sale_lag_1'),
(0.0, 'item_type_id_cnt_month_mean_Expanding_lag_4'),
(0.0, 'item_subtype_id_cnt_month_mean_Expanding_lag_2'),
(0.0, 'item_subtype_id_cnt_month_mean_Expanding_lag_12'),
(0.0, 'item_subtype_id_cnt_month_mean_Expanding_lag_1'),
(0.0, 'item_category_id_cnt_month_mean_Expanding_lag_4'),
(0.0, 'item_category_id_cnt_month_mean_Expanding_lag_3'),
(0.0, 'item_category_id_cnt_month_mean_Expanding_lag_2'),
(0.0, 'item_category_id_cnt_month_mean_Expanding_lag_12'),
(0.0, 'delta_revenue_lag_1'),
(0.0, 'city_id_cnt_month_mean_Expanding_lag_4'),
(0.0, 'city_id_cnt_month_mean_Expanding_lag_3'),
(0.0, 'city_id_cnt_month_mean_Expanding_lag_2'),
(0.0, 'city_id_cnt_month_mean_Expanding_lag_12'),
(0.0, 'city_id_cnt_month_mean_Expanding_lag_1'),
(0.0, 'cat_subtype_id_cnt_month_mean_Expanding_lag_2'),
(0.0, 'cat_subtype_id_cnt_month_mean_Expanding_lag_1')]
```

```
[120]: x_lst = x_lst[:30]
features_x = []
for i in x_lst:
    features_x.append(i[1])
```

```
[121]: pd.set_option('display.max_rows', None)
# put all selection together
feature_selection_df = pd.DataFrame({'Feature': feature_name, 'Pearson':
    ↳ cor_support, 'Chi-2': chi_support, 'RFE': rfe_support})
```

```
[122]: rf_df = pd.DataFrame({'Feature': features_rf, 'RF': True})
gbm_df = pd.DataFrame({'Feature': features_gbm, 'LGBM': True})
xgb_df = pd.DataFrame({'Feature': features_x, 'XGB': True})
```

```
[123]: feature_selection_df = pd.merge(feature_selection_df, rf_df, how='left',
    ↳ on='Feature')
feature_selection_df = pd.merge(feature_selection_df, gbm_df, how='left',
    ↳ on='Feature')
feature_selection_df = pd.merge(feature_selection_df, xgb_df, how='left',
    ↳ on='Feature')
feature_selection_df = feature_selection_df.fillna(False)
```

```
[124]: # count the selected times for each feature
feature_selection_df['Total'] = np.sum(feature_selection_df, axis=1)
# display the top
feature_selection_df = feature_selection_df.sort_values(['Total', 'Feature'],
→ascending=False)
feature_selection_df.index = range(1, len(feature_selection_df)+1)
feature_selection_df
```

```
[124]:
```

	Feature	Pearson	Chi-2	RFE \
1	shop_item_id_cnt_month_mean_Smooth_lag_2	True	True	True
2	shop_item_id_cnt_month_mean_Smooth_lag_1	True	True	True
3	item_cnt_month_lag_2	True	True	True
4	item_cnt_month_lag_1	True	True	True
5	shop_item_id_cnt_month_mean_Smooth_lag_4	True	True	True
6	months_since_shop_item_first_sale_lag_1	True	True	False
7	item_id_cnt_month_mean_Expanding_lag_4	True	True	False
8	item_id_cnt_month_mean_Expanding_lag_2	True	True	True
9	item_id_cnt_month_mean_Expanding_lag_1	True	True	True
10	item_cnt_month_lag_3	True	True	True
11	shop_item_id_cnt_month_mean_Smooth_lag_3	True	True	True
12	shop_item_id_cnt_month_mean_Smooth_lag_12	True	False	True
13	item_type_id_cnt_month_mean_Expanding_lag_2	False	False	True
14	item_subtype_id_cnt_month_mean_Expanding_lag_2	True	False	True
15	item_subtype_id_cnt_month_mean_Expanding_lag_1	True	False	True
16	item_cnt_month_lag_6	True	True	False
17	months_since_item_shop_last_sale_lag_1	False	False	False
18	months_since_item_first_sale_lag_1	False	True	False
19	month	False	False	False
20	item_type_id_cnt_month_mean_Expanding_lag_1	False	False	True
21	item_id_cnt_month_mean_Expanding_lag_12	False	True	False
22	item_cnt_month_lag_8	True	True	False
23	item_cnt_month_lag_5	True	True	False
24	item_category_id_cnt_month_mean_Expanding_lag_1	True	False	True
25	item_category_id	False	False	False
26	cat_type_id_cnt_month_mean_Expanding_lag_1	False	False	True
27	shop_type_id_cnt_month_mean_Expanding_lag_1	False	False	True
28	shop_id_cnt_month_mean_Expanding_lag_4	False	True	False
29	item_type_id	False	False	False
30	item_subtype_id_cnt_month_mean_Expanding_lag_4	False	False	True
31	item_subtype_id	False	False	False
32	item_id_cnt_month_mean_Expanding_lag_3	True	True	False
33	item_id	False	False	False
34	item_cnt_month_lag_9	True	True	False
35	item_cnt_month_lag_7	True	True	False
36	item_cnt_month_lag_4	True	True	False
37	item_cnt_month_lag_12	True	True	False
38	item_cnt_month_lag_11	True	True	False

39	item_cnt_month_lag_10	True	True	False
40	item_category_id_cnt_month_mean_Expanding_lag_3	True	False	True
41	item_category_id_cnt_month_mean_Expanding_lag_2	True	False	True
42	delta_price_lag	False	False	False
43	days_of_month	False	False	False
44	city_id_cnt_month_mean_Expanding_lag_2	False	True	False
45	city_id_cnt_month_mean_Expanding_lag_1	False	True	False
46	city_id	False	False	False
47	cat_subtype_id_cnt_month_mean_Expanding_lag_4	False	False	True
48	cat_subtype_id_cnt_month_mean_Expanding_lag_3	True	False	True
49	cat_subtype_id_cnt_month_mean_Expanding_lag_2	True	False	True
50	cat_subtype_id_cnt_month_mean_Expanding_lag_1	True	False	True
51	cat_subtype_id	False	False	False
52	year	False	False	False
53	shop_type_id_cnt_month_mean_Expanding_lag_4	False	False	False
54	shop_type_id_cnt_month_mean_Expanding_lag_3	False	False	False
55	shop_type_id_cnt_month_mean_Expanding_lag_12	False	False	True
56	shop_type_id	False	False	False
57	shop_id_cnt_month_mean_Expanding_lag_3	False	True	False
58	shop_id_cnt_month_mean_Expanding_lag_2	False	True	False
59	shop_id_cnt_month_mean_Expanding_lag_12	False	False	False
60	shop_id_cnt_month_mean_Expanding_lag_1	False	True	False
61	shop_id	False	False	False
62	months_since_item_last_sale_lag_1	False	False	True
63	item_subtype_id_cnt_month_mean_Expanding_lag_3	False	False	True
64	item_category_id_cnt_month_mean_Expanding_lag_4	False	False	True
65	item_category_id_cnt_month_mean_Expanding_lag_12	False	False	True
66	delta_revenue_lag_1	False	False	False
67	city_id_cnt_month_mean_Expanding_lag_3	False	True	False
68	cat_type_id_cnt_month_mean_Expanding_lag_4	False	False	False
69	cat_type_id_cnt_month_mean_Expanding_lag_3	False	False	False
70	cat_type_id_cnt_month_mean_Expanding_lag_2	False	False	False
71	cat_type_id_cnt_month_mean_Expanding_lag_12	False	False	False
72	cat_type_id	False	False	False
73	cat_subtype_id_cnt_month_mean_Expanding_lag_12	False	False	True
74	shop_type_id_cnt_month_mean_Expanding_lag_2	False	False	False
75	item_type_id_cnt_month_mean_Expanding_lag_4	False	False	False
76	item_type_id_cnt_month_mean_Expanding_lag_3	False	False	False
77	item_type_id_cnt_month_mean_Expanding_lag_12	False	False	False
78	item_subtype_id_cnt_month_mean_Expanding_lag_12	False	False	False
79	city_id_cnt_month_mean_Expanding_lag_4	False	False	False
80	city_id_cnt_month_mean_Expanding_lag_12	False	False	False

	RF	LGBM	XGB	Total
1	True	True	True	6
2	True	True	True	6
3	True	True	True	6

4	True	True	True	6
5	False	True	True	5
6	True	True	True	5
7	True	True	True	5
8	True	False	True	5
9	True	True	False	5
10	False	True	True	5
11	False	False	True	4
12	False	True	True	4
13	True	True	True	4
14	True	True	False	4
15	True	True	False	4
16	False	True	True	4
17	True	True	True	3
18	True	False	True	3
19	True	True	True	3
20	True	True	False	3
21	False	True	True	3
22	False	False	True	3
23	False	False	True	3
24	False	True	False	3
25	True	True	True	3
26	True	True	False	3
27	True	False	False	2
28	False	False	True	2
29	True	True	False	2
30	False	False	True	2
31	True	True	False	2
32	False	False	False	2
33	True	True	False	2
34	False	False	False	2
35	False	False	False	2
36	False	False	False	2
37	False	False	False	2
38	False	False	False	2
39	False	False	False	2
40	False	False	False	2
41	False	False	False	2
42	True	True	False	2
43	True	True	False	2
44	True	False	False	2
45	True	False	False	2
46	True	False	True	2
47	False	False	True	2
48	False	False	False	2
49	False	False	False	2
50	False	False	False	2

51	True	True	False	2
52	False	True	False	1
53	False	False	True	1
54	False	False	True	1
55	False	False	False	1
56	True	False	False	1
57	False	False	False	1
58	False	False	False	1
59	False	False	True	1
60	False	False	False	1
61	True	False	False	1
62	False	False	False	1
63	False	False	False	1
64	False	False	False	1
65	False	False	False	1
66	True	False	False	1
67	False	False	False	1
68	False	False	True	1
69	False	False	True	1
70	False	True	False	1
71	False	False	True	1
72	False	True	False	1
73	False	False	False	1
74	False	False	False	0
75	False	False	False	0
76	False	False	False	0
77	False	False	False	0
78	False	False	False	0
79	False	False	False	0
80	False	False	False	0

```
[128]: #select from sales df features that will be used for modelling and drop
↳ unwanted ones
unwanted = feature_selection_df.Feature[74:]
```

```
[129]: to_drop = []
for i in unwanted:
    to_drop.append(i)
```

```
[130]: sales_df = sales.drop(to_drop, axis=1)
```

```
[131]: sales_df.head()
```

```
[131]:
```

	cat_subtype_id	cat_type_id	city_id	date_block_num	\
3395293	0.062500	0.555556	0.0	12	
3395294	0.015625	0.555556	0.0	12	
3395295	0.062500	0.555556	0.0	12	

3395296	0.015625	0.555556	0.0	12
3395297	0.062500	0.555556	0.0	12

	item_category_id	item_cnt_month	item_id	item_subtype_id	\
3395293	0.469136	0.0	0.001308	0.025225	
3395294	0.432099	0.0	0.001353	0.337538	
3395295	0.469136	1.0	0.001398	0.025225	
3395296	0.432099	1.0	0.001444	0.337538	
3395297	0.469136	0.0	0.001489	0.821021	

	item_type_id	shop_id	shop_type_id	item_cnt_month_lag_1	\
3395293	0.022989	0.0	1.0	0.00	
3395294	0.022989	0.0	1.0	0.00	
3395295	0.022989	0.0	1.0	0.00	
3395296	0.022989	0.0	1.0	0.05	
3395297	0.022989	0.0	1.0	0.00	

	item_cnt_month_lag_2	item_cnt_month_lag_3	item_cnt_month_lag_4	\
3395293	0.0	0.0	0.0	
3395294	0.0	0.0	0.0	
3395295	0.0	0.0	0.0	
3395296	0.1	0.0	0.0	
3395297	0.0	0.0	0.0	

	item_cnt_month_lag_5	item_cnt_month_lag_6	item_cnt_month_lag_7	\
3395293	0.0	0.0	0.05	
3395294	0.0	0.0	0.00	
3395295	0.0	0.0	0.00	
3395296	0.0	0.0	0.00	
3395297	0.0	0.0	0.00	

	item_cnt_month_lag_8	item_cnt_month_lag_9	item_cnt_month_lag_10	\
3395293	0.0	0.00	0.05	
3395294	0.0	0.05	0.05	
3395295	0.0	0.00	0.00	
3395296	0.0	0.00	0.00	
3395297	0.0	0.00	0.00	

	item_cnt_month_lag_11	item_cnt_month_lag_12	\
3395293	0.0	0.00	
3395294	0.2	0.00	
3395295	0.0	0.00	
3395296	0.0	0.05	
3395297	0.0	0.00	

	shop_id_cnt_month_mean_Expanding_lag_1	\
3395293	0.081879	

3395294	0.081878
3395295	0.081877
3395296	0.081877
3395297	0.081885
shop_id_cnt_month_mean_Expanding_lag_2 ... \	
3395293	0.081114 ...
3395294	0.081113 ...
3395295	0.081112 ...
3395296	0.081111 ...
3395297	0.081131 ...
item_type_id_cnt_month_mean_Expanding_lag_1 \	
3395293	0.012491
3395294	0.012491
3395295	0.012491
3395296	0.012491
3395297	0.012491
item_type_id_cnt_month_mean_Expanding_lag_2 \	
3395293	0.012046
3395294	0.012046
3395295	0.012046
3395296	0.012046
3395297	0.012046
item_subtype_id_cnt_month_mean_Expanding_lag_1 \	
3395293	0.020679
3395294	0.010859
3395295	0.020679
3395296	0.010859
3395297	0.010343
item_subtype_id_cnt_month_mean_Expanding_lag_2 \	
3395293	0.019118
3395294	0.010219
3395295	0.019118
3395296	0.010219
3395297	0.009702
item_subtype_id_cnt_month_mean_Expanding_lag_3 \	
3395293	0.019113
3395294	0.010387
3395295	0.019113
3395296	0.010387
3395297	0.009787

	item_subtype_id_cnt_month_mean_Expanding_lag_4 \
3395293	0.018981
3395294	0.010396
3395295	0.018981
3395296	0.010396
3395297	0.009994
	cat_type_id_cnt_month_mean_Expanding_lag_1 \
3395293	0.037261
3395294	0.037261
3395295	0.037261
3395296	0.037261
3395297	0.037261
	cat_type_id_cnt_month_mean_Expanding_lag_2 \
3395293	0.037322
3395294	0.037322
3395295	0.037322
3395296	0.037322
3395297	0.037322
	cat_type_id_cnt_month_mean_Expanding_lag_3 \
3395293	0.03756
3395294	0.03756
3395295	0.03756
3395296	0.03756
3395297	0.03756
	cat_type_id_cnt_month_mean_Expanding_lag_4 \
3395293	0.037923
3395294	0.037923
3395295	0.037923
3395296	0.037923
3395297	0.037922
	cat_type_id_cnt_month_mean_Expanding_lag_12 \
3395293	0.000000
3395294	0.000000
3395295	0.000000
3395296	0.000000
3395297	0.041667
	cat_subtype_id_cnt_month_mean_Expanding_lag_1 \
3395293	0.016609
3395294	0.013401
3395295	0.016609
3395296	0.013401

3395297	0.016609			
	cat_subtype_id_cnt_month_mean_Expanding_lag_2 \			
3395293	0.016573			
3395294	0.013517			
3395295	0.016573			
3395296	0.013517			
3395297	0.016573			
	cat_subtype_id_cnt_month_mean_Expanding_lag_3 \			
3395293	0.016583			
3395294	0.013746			
3395295	0.016583			
3395296	0.013746			
3395297	0.016583			
	cat_subtype_id_cnt_month_mean_Expanding_lag_4 \			
3395293	0.016722			
3395294	0.013844			
3395295	0.016722			
3395296	0.013844			
3395297	0.016722			
	cat_subtype_id_cnt_month_mean_Expanding_lag_12	delta_price_lag \		
3395293	0.000000	0.130395		
3395294	0.000000	0.218334		
3395295	0.000000	0.148123		
3395296	0.020995	0.195501		
3395297	0.000000	0.251542		
	delta_revenue_lag_1	month	year	days_of_month \
3395293	0.425059	0.0	0.0	1.0
3395294	0.425059	0.0	0.0	1.0
3395295	0.425059	0.0	0.0	1.0
3395296	0.425059	0.0	0.0	1.0
3395297	0.425059	0.0	0.0	1.0
	months_since_item_shop_last_sale_lag_1 \			
3395293	0.0625			
3395294	0.0625			
3395295	0.0000			
3395296	0.0625			
3395297	0.0000			
	months_since_shop_item_first_sale_lag_1 \			
3395293	0.000000			
3395294	0.000000			

3395295	0.000000
3395296	0.363636
3395297	0.000000

	months_since_item_last_sale_lag_1	months_since_item_first_sale_lag_1
3395293	0.0625	0.000000
3395294	0.0625	0.000000
3395295	0.0625	0.000000
3395296	0.0625	0.363636
3395297	0.0625	0.000000

[5 rows x 76 columns]

```
[132]: #save final dataframe
sales_df.to_pickle('data_6.pickle.gzde', compression='gzip')
```

```
[133]: del sales_df
del sales
```

1.3 Modelling

```
[3]: sales = pd.read_pickle('data_6.pickle.gzde', compression='gzip')
```

```
[18]: #first-level model
#save date_block_num, as cannot use as feature, but need it to split dataset
↳ into parts
feature_cols = list(sales)
feature_cols = [e for e in feature_cols if e not in
↳ ('date_block_num', 'item_cnt_month')]
num_first_level_models = 5

dates = sales['date_block_num']
last_block = dates.max()
print('Test `date_block_num` is %d' % last_block)
print(feature_cols)
```

```
Test `date_block_num` is 34
['cat_subtype_id', 'cat_type_id', 'city_id', 'item_category_id', 'item_id',
'item_subtype_id', 'item_type_id', 'shop_id', 'shop_type_id',
'item_cnt_month_lag_1', 'item_cnt_month_lag_2', 'item_cnt_month_lag_3',
'item_cnt_month_lag_4', 'item_cnt_month_lag_5', 'item_cnt_month_lag_6',
'item_cnt_month_lag_7', 'item_cnt_month_lag_8', 'item_cnt_month_lag_9',
'item_cnt_month_lag_10', 'item_cnt_month_lag_11', 'item_cnt_month_lag_12',
'shop_id_cnt_month_mean_Expanding_lag_1',
'shop_id_cnt_month_mean_Expanding_lag_2',
'shop_id_cnt_month_mean_Expanding_lag_3',
'shop_id_cnt_month_mean_Expanding_lag_4',
```

```

'shop_id_cnt_month_mean_Expanding_lag_12',
'item_id_cnt_month_mean_Expanding_lag_1',
'item_id_cnt_month_mean_Expanding_lag_2',
'item_id_cnt_month_mean_Expanding_lag_3',
'item_id_cnt_month_mean_Expanding_lag_4',
'item_id_cnt_month_mean_Expanding_lag_12',
'shop_item_id_cnt_month_mean_Smooth_lag_1',
'shop_item_id_cnt_month_mean_Smooth_lag_2',
'shop_item_id_cnt_month_mean_Smooth_lag_3',
'shop_item_id_cnt_month_mean_Smooth_lag_4',
'shop_item_id_cnt_month_mean_Smooth_lag_12',
'shop_type_id_cnt_month_mean_Expanding_lag_1',
'shop_type_id_cnt_month_mean_Expanding_lag_2',
'shop_type_id_cnt_month_mean_Expanding_lag_3',
'shop_type_id_cnt_month_mean_Expanding_lag_4',
'shop_type_id_cnt_month_mean_Expanding_lag_12',
'city_id_cnt_month_mean_Expanding_lag_1',
'city_id_cnt_month_mean_Expanding_lag_2',
'city_id_cnt_month_mean_Expanding_lag_3',
'item_category_id_cnt_month_mean_Expanding_lag_1',
'item_category_id_cnt_month_mean_Expanding_lag_2',
'item_category_id_cnt_month_mean_Expanding_lag_3',
'item_category_id_cnt_month_mean_Expanding_lag_4',
'item_category_id_cnt_month_mean_Expanding_lag_12',
'item_type_id_cnt_month_mean_Expanding_lag_1',
'item_type_id_cnt_month_mean_Expanding_lag_2',
'item_subtype_id_cnt_month_mean_Expanding_lag_1',
'item_subtype_id_cnt_month_mean_Expanding_lag_2',
'item_subtype_id_cnt_month_mean_Expanding_lag_3',
'item_subtype_id_cnt_month_mean_Expanding_lag_4',
'cat_type_id_cnt_month_mean_Expanding_lag_1',
'cat_type_id_cnt_month_mean_Expanding_lag_2',
'cat_type_id_cnt_month_mean_Expanding_lag_3',
'cat_type_id_cnt_month_mean_Expanding_lag_4',
'cat_type_id_cnt_month_mean_Expanding_lag_12',
'cat_subtype_id_cnt_month_mean_Expanding_lag_1',
'cat_subtype_id_cnt_month_mean_Expanding_lag_2',
'cat_subtype_id_cnt_month_mean_Expanding_lag_3',
'cat_subtype_id_cnt_month_mean_Expanding_lag_4',
'cat_subtype_id_cnt_month_mean_Expanding_lag_12', 'delta_price_lag',
'delta_revenue_lag_1', 'month', 'year', 'days_of_month',
'months_since_item_shop_last_sale_lag_1',
'months_since_shop_item_first_sale_lag_1', 'months_since_item_last_sale_lag_1',
'months_since_item_first_sale_lag_1']

```

```
[19]: start_first_level_total = time.perf_counter()
```

```
scoringMethod = 'r2'; from sklearn.metrics import mean_squared_error; from math_
↳import sqrt
```

```
[20]: #train meta-features

months_to_generate_meta_features = range(24,last_block + 1)
mask = dates.isin(months_to_generate_meta_features)
Target = 'item_cnt_month'
y_all_level2 = sales[Target][mask].values
X_all_level2 = np.zeros([y_all_level2.shape[0], num_first_level_models])
```

Parameters tuned on separate notebook. The optimal parameters found are as follows:

XGBoost:

```
{'colsample_bytree': 0.6,      'eta': 0.3,      'max_depth': 10,
'min_child_weight': 300,      'seed': 42,      'subsample': 1.0,
'n_estimators': 125}
```

Light GBM:

```
{'bagging_fraction': 0.2,      'feature_fraction': 0.8,      'learning_rate':
0.05,      'max_depth': -1,      'min_data_in_leaf': 300,      'num_leaves': 128,
'seed': 42,      'n_estimators': 200}
```

Random Forest:

```
{'bootstrap': True,      'max_depth': 20,      'max_features': 'sqrt',
'min_samples_leaf': 4,      'min_samples_split': 5,      'n_estimators': 200}
```

SGD:

```
{'alpha': 0.0001,      'penalty': 'l1'}
```

Keras:

```
{'activation': 'tanh',      'activation2': 'relu',      'batch_size': 5000,
'dropout_rate': 0.0,      'epochs': 50,      'init': 'he_normal',      'neurons':
200,      'optimizer': 'Adam'}
```

```
[21]: #hyperparameters tuned on separate notebook and loaded here
#xgb params
with open("XBG_Params.pkl", 'rb') as file:
    xgb_params = pickle.load(file)

xgb_params
```

```
[21]: {'colsample_bytree': 0.6,
'eta': 0.3,
'max_depth': 10,
'min_child_weight': 300,
'seed': 42,
'subsample': 1.0,
'n_estimators': 125}
```

```
[22]: with open("LGB_Params.pkl", 'rb') as file:
      lgb_params = pickle.load(file)

lgb_params
```

```
[22]: {'bagging_fraction': 0.2,
      'feature_fraction': 0.8,
      'learning_rate': 0.05,
      'max_depth': -1,
      'min_data_in_leaf': 300,
      'num_leaves': 128,
      'seed': 42,
      'n_estimators': 200}
```

```
[23]: with open("RF_Params.pkl", 'rb') as file:
      rf_params = pickle.load(file)

rf_params
```

```
[23]: {'bootstrap': True,
      'max_depth': 20,
      'max_features': 'sqrt',
      'min_samples_leaf': 4,
      'min_samples_split': 5,
      'n_estimators': 200}
```

```
[24]: with open("SGD_Params.pkl", 'rb') as file:
      sgd_params = pickle.load(file)

sgd_params
```

```
[24]: {'alpha': 0.0001, 'penalty': 'l1', 'random_state': 0}
```

```
[25]: with open("Keras_Params.pkl", 'rb') as file:
      keras_params = pickle.load(file)

keras_params
```

```
[25]: {'activation': 'tanh',
      'activation2': 'relu',
      'batch_size': 5000,
      'dropout_rate': 0.0,
      'epochs': 50,
      'init': 'he_normal',
      'neurons': 200,
      'optimizer': 'Adam'}
```

```

[26]: #fill `X_train_level2` with metafeatures
slice_start = 0

for cur_block_num in tqdm(months_to_generate_meta_features):
    print('-' * 50)
    print('Start training for month%d'% cur_block_num)
    start_cur_month = time.perf_counter()

    cur_X_train = sales.loc[dates < cur_block_num][feature_cols]
    cur_X_test = sales.loc[dates == cur_block_num][feature_cols]

    cur_y_train = sales.loc[dates < cur_block_num, Target].values
    cur_y_test = sales.loc[dates == cur_block_num, Target].values

    # Create Numpy arrays of train, test and target dataframes to feed into
    ↪models
    train_x = cur_X_train.values
    train_y = cur_y_train.ravel()
    test_x = cur_X_test.values
    test_y = cur_y_test.ravel()

    preds = []

    from sklearn.linear_model import (LinearRegression, SGDRegressor)
    import lightgbm as lgb
    import xgboost as xgb
    from sklearn.ensemble import RandomForestRegressor

    sgdr = SGDRegressor(penalty = 'l1', alpha = 0.0001, random_state = SEED )

    rf = RandomForestRegressor(
        bootstrap=True,
        max_depth=20,
        max_features='sqrt',
        min_samples_leaf=4,
        min_samples_split=5,
        n_estimators=200)

    lgb_params = {
        'feature_fraction': 0.8,
        'metric': 'rmse',
        'min_data_in_leaf': 300,
        'bagging_fraction': 0.2,
        'learning_rate': 0.05,
        'max_depth': -1,
        'objective': 'mse',
        'num_leaves': 128,

```

```

        'seed': 42,
        'n_estimators': 200,
        'verbose': 0
    }

    xgb_params = {
        'colsample_bytree': 0.6,
        'eta': 0.3,
        'max_depth': 10,
        'min_child_weight': 300,
        'n_estimators': 125,
        'seed': 42,
        'subsample': 1.0,
        'verbosity': 0
    }

    print('Training Model %d: %s'%(len(preds), 'XGBoost'))

    start = time.perf_counter()

    estimator = xgb.train(xgb_params, xgb.DMatrix(data=train_x, label=train_y))
    pred_test = estimator.predict(xgb.DMatrix(test_x))
    preds.append(pred_test)

    run = time.perf_counter() - start

    print('{} runs for {:.2f} seconds.'.format('XGBoost', run))
    print()

    print('Training Model %d: %s'%(len(preds), 'Lightgbm'))

    start = time.perf_counter()

    estimator = lgb.train(lgb_params, lgb.Dataset(train_x, label=train_y))
    pred_test = estimator.predict(test_x)
    preds.append(pred_test)

    run = time.perf_counter() - start

    print('{} runs for {:.2f} seconds.'.format('Lightgbm', run))
    print()

    estimators = [sgdr, rf]
    for estimator in estimators:
        print('Training Model %d: %s'%(len(preds), estimator.__class__.__name__))
        start = time.perf_counter()

```



```

    estimator.fit(train_x, train_y)
    pred_test = estimator.predict(test_x)
    preds.append(pred_test)

    run = time.perf_counter() - start
    print('{} runs for {:.2f} seconds.'.format(estimator.__class__.__name__, run))
    print()

    print('Training Model %d: %s'%(len(preds), 'Keras'))

    start = time.perf_counter()

    from keras.models import Sequential
    from keras.layers import Dense
    from keras.wrappers.scikit_learn import KerasRegressor

    def baseline_model():
        # create model
        model = Sequential()
        model.add(Dense(200, input_dim=train_x.shape[1],
        ↪kernel_initializer='he_normal', activation='tanh'))
        model.add(Dense(1, kernel_initializer='he_normal', activation='relu'))
        # Compile model
        model.compile(loss='mse', optimizer='Adam', metrics=['mse'])
        return model

    estimator = KerasRegressor(build_fn=baseline_model, verbose=0, epochs=50,
    ↪batch_size = 5000)

    estimator.fit(train_x, train_y)
    pred_test = estimator.predict(test_x)
    preds.append(pred_test)

    run = time.perf_counter() - start
    print('{} runs for {:.2f} seconds.'.format('Keras', run))

    cur_month_run_total = time.perf_counter() - start_cur_month

    print('Total running time was {:.2f} minutes.'.format(cur_month_run_total/
    ↪60))
    print('-' * 50)

    slice_end = slice_start + cur_X_test.shape[0]

```

```
X_all_level2[ slice_start : slice_end , :] = np.c_[preds].transpose()  
slice_start = slice_end
```

```
0%|          | 0/11 [00:00<?, ?it/s]
```

```
-----  
Start training for month24  
Training Model 0: XGBoost  
XGBoost runs for 46.14 seconds.
```

```
Training Model 1: Lightgbm  
Lightgbm runs for 57.14 seconds.
```

```
Training Model 2: SGDRegressor  
SGDRegressor runs for 16.68 seconds.
```

```
Training Model 3: RandomForestRegressor  
RandomForestRegressor runs for 3388.77 seconds.
```

```
Training Model 4: Keras
```

```
9%|          | 1/11 [1:07:05<11:10:54, 4025.46s/it]
```

```
Keras runs for 515.55 seconds.  
Total running time was 67.09 minutes.  
-----  
-----
```

```
Start training for month25  
Training Model 0: XGBoost  
XGBoost runs for 52.51 seconds.
```

```
Training Model 1: Lightgbm  
Lightgbm runs for 59.99 seconds.
```

```
Training Model 2: SGDRegressor  
SGDRegressor runs for 18.36 seconds.
```

```
Training Model 3: RandomForestRegressor  
RandomForestRegressor runs for 4280.31 seconds.
```

```
Training Model 4: Keras
```

```
18%|         | 2/11 [3:37:06<13:47:43, 5518.18s/it]
```

Keras runs for 523.81 seconds.
Total running time was 82.27 minutes.

Start training for month26
Training Model 0: XGBoost
XGBoost runs for 53.21 seconds.

Training Model 1: Lightgbm
Lightgbm runs for 59.85 seconds.

Training Model 2: SGDRegressor
SGDRegressor runs for 18.45 seconds.

Training Model 3: RandomForestRegressor
RandomForestRegressor runs for 3807.67 seconds.

Training Model 4: Keras

27%| | 3/11 [4:52:04<11:34:55, 5211.99s/it]

Keras runs for 556.77 seconds.
Total running time was 74.96 minutes.

Start training for month27
Training Model 0: XGBoost
XGBoost runs for 53.33 seconds.

Training Model 1: Lightgbm
Lightgbm runs for 55.10 seconds.

Training Model 2: SGDRegressor
SGDRegressor runs for 21.02 seconds.

Training Model 3: RandomForestRegressor
RandomForestRegressor runs for 4319.26 seconds.

Training Model 4: Keras

36%| | 4/11 [6:16:06<10:02:08, 5161.17s/it]

Keras runs for 592.12 seconds.
Total running time was 84.04 minutes.

Start training for month28
Training Model 0: XGBoost
XGBoost runs for 55.97 seconds.

Training Model 1: Lightgbm
Lightgbm runs for 56.19 seconds.

Training Model 2: SGDRegressor
SGDRegressor runs for 22.67 seconds.

Training Model 3: RandomForestRegressor
RandomForestRegressor runs for 4601.93 seconds.

Training Model 4: Keras

45%| | 5/11 [7:45:08<8:41:32, 5215.36s/it]

Keras runs for 603.41 seconds.
Total running time was 89.03 minutes.

Start training for month29
Training Model 0: XGBoost
XGBoost runs for 60.19 seconds.

Training Model 1: Lightgbm
Lightgbm runs for 59.14 seconds.

Training Model 2: SGDRegressor
SGDRegressor runs for 21.60 seconds.

Training Model 3: RandomForestRegressor
RandomForestRegressor runs for 4891.86 seconds.

Training Model 4: Keras

55%| | 6/11 [9:19:45<7:26:09, 5353.92s/it]

Keras runs for 642.83 seconds.
Total running time was 94.62 minutes.

Start training for month30
Training Model 0: XGBoost
XGBoost runs for 82.18 seconds.

Training Model 1: Lightgbm
Lightgbm runs for 62.91 seconds.

Training Model 2: SGDRegressor
SGDRegressor runs for 22.56 seconds.

Training Model 3: RandomForestRegressor
RandomForestRegressor runs for 5224.10 seconds.

Training Model 4: Keras

64%| | 7/11 [11:01:01<6:11:22, 5570.53s/it]

Keras runs for 682.56 seconds.
Total running time was 101.27 minutes.

Start training for month31
Training Model 0: XGBoost
XGBoost runs for 85.40 seconds.

Training Model 1: Lightgbm
Lightgbm runs for 65.56 seconds.

Training Model 2: SGDRegressor
SGDRegressor runs for 23.92 seconds.

Training Model 3: RandomForestRegressor
RandomForestRegressor runs for 5535.80 seconds.

Training Model 4: Keras

73%| | 8/11 [12:47:59<4:51:14, 5824.73s/it]

Keras runs for 705.44 seconds.
Total running time was 106.96 minutes.

Start training for month32
Training Model 0: XGBoost
XGBoost runs for 91.65 seconds.

Training Model 1: Lightgbm
Lightgbm runs for 68.30 seconds.

Training Model 2: SGDRegressor

SGDRegressor runs for 24.95 seconds.

Training Model 3: RandomForestRegressor
RandomForestRegressor runs for 5797.19 seconds.

Training Model 4: Keras

82%| | 9/11 [14:40:07<3:23:11, 6095.80s/it]

Keras runs for 744.41 seconds.
Total running time was 112.14 minutes.

Start training for month33
Training Model 0: XGBoost
XGBoost runs for 96.42 seconds.

Training Model 1: Lightgbm
Lightgbm runs for 71.00 seconds.

Training Model 2: SGDRegressor
SGDRegressor runs for 26.53 seconds.

Training Model 3: RandomForestRegressor
RandomForestRegressor runs for 6080.74 seconds.

Training Model 4: Keras

91%| | 10/11 [16:37:54<1:46:26, 6386.90s/it]

Keras runs for 789.55 seconds.
Total running time was 117.77 minutes.

Start training for month34
Training Model 0: XGBoost
XGBoost runs for 101.11 seconds.

Training Model 1: Lightgbm
Lightgbm runs for 74.35 seconds.

Training Model 2: SGDRegressor
SGDRegressor runs for 27.23 seconds.

Training Model 3: RandomForestRegressor
RandomForestRegressor runs for 6312.06 seconds.

Training Model 4: Keras

100%| | 11/11 [19:06:30<00:00, 6253.64s/it]

Keras runs for 884.63 seconds.

Total running time was 123.36 minutes.

[27]: *#split train and test*

```
test_nrow = len(preds[0])
X_train_level2 = X_all_level2[ : -test_nrow, :]
X_test_level2 = X_all_level2[ -test_nrow: , :]
y_train_level2 = y_all_level2[ : -test_nrow]
y_test_level2 = y_all_level2[ -test_nrow : ]
```

[28]: *#RMSE for individual models*

```
xgb_pred = X_train_level2[:, 0].clip(0, 20)
lgb_pred = X_train_level2[:, 1].clip(0, 20)
sgdr_pred = X_train_level2[:, 2].clip(0, 20)
rf_pred = X_train_level2[:, 3].clip(0, 20)
keras_pred = X_train_level2[:, 4].clip(0, 20)
print('Train RMSE for %s is %f' %('xgb_pred',
    ↳sqrt(mean_squared_error(y_train_level2, xgb_pred))))
print('Train RMSE for %s is %f' %('lgb_pred',
    ↳sqrt(mean_squared_error(y_train_level2, lgb_pred))))
print('Train RMSE for %s is %f' %('sgdr_pred',
    ↳sqrt(mean_squared_error(y_train_level2, sgdr_pred))))
print('Train RMSE for %s is %f' %('rf_pred',
    ↳sqrt(mean_squared_error(y_train_level2, rf_pred))))
print('Train RMSE for %s is %f' %('keras_pred',
    ↳sqrt(mean_squared_error(y_train_level2, keras_pred))))
```

Train RMSE for xgb_pred is 0.767315

Train RMSE for lgb_pred is 0.749203

Train RMSE for sgdr_pred is 0.856755

Train RMSE for rf_pred is 0.760995

Train RMSE for keras_pred is 0.866221

[29]: *#Ensembling*

```
pred_list = {}
```

```
#second level learning model via linear regression
```

```
print('Training Second level learning model via linear regression')
```

```

from sklearn.linear_model import (LinearRegression, SGDRegressor)

lr = LinearRegression()
lr.fit(X_train_level2, y_train_level2)

#compute R-squared on train and test sets
# print('Train R-squared for %s is %f' %('test_preds_lr_stacking',
    →sqrt(mean_squared_error(y_train_level2, lr.predict(X_train_level2)))))

test_preds_lr_stacking = lr.predict(X_test_level2)
train_preds_lr_stacking = lr.predict(X_train_level2)
print('Train RMSE for %s is %f' %('train_preds_lr_stacking',
    →sqrt(mean_squared_error(y_train_level2, train_preds_lr_stacking))))

pred_list['test_preds_lr_stacking'] = test_preds_lr_stacking
if Validation:
    print('Test RMSE for %s is %f' %('test_preds_lr_stacking',
    →sqrt(mean_squared_error(y_test_level2, test_preds_lr_stacking))))

```

Training Second level learning model via linear regression
 Train RMSE for train_preds_lr_stacking is 0.745656

```

[30]: #second level learning model via SGDRegressor
print('Training Second level learning model via SGDRegressor')
sgdr = SGDRegressor(
    penalty = 'l2' ,
    random_state = SEED )

sgdr.fit(X_train_level2, y_train_level2)

#compute R-squared on train and test sets
# print('Train R-squared for %s is %f' %('test_preds_lr_stacking',
    →sqrt(mean_squared_error(y_train_level2, lr.predict(X_train_level2)))))
test_preds_sgdr_stacking = sgdr.predict(X_test_level2)
train_preds_sgdr_stacking = sgdr.predict(X_train_level2)
print('Train RMSE for %s is %f' %('train_preds_lr_stacking',
    →sqrt(mean_squared_error(y_train_level2, train_preds_sgdr_stacking))))

pred_list['test_preds_sgdr_stacking'] = test_preds_sgdr_stacking

if Validation:
    print('Test RMSE for %s is %f' %('test_preds_sgdr_stacking',
    →sqrt(mean_squared_error(y_test_level2, test_preds_sgdr_stacking))))

#print('%0.2f min: Finish training second level model'%((time.time() -
    →start_time)/60))

```


Training Second level learning model via SGDRegressor
Train RMSE for train_preds_lr_stacking is 0.761842

```
[31]: #submission
if not Validation:
    submission = pd.read_csv('sample_submission.csv')
    ver = 3

    for pred_ver in ['lr_stacking', 'sgdr_stacking']:
        print(pred_list['test_preds_' + pred_ver].clip(0,20).mean())
        submission['item_cnt_month'] = pred_list['test_preds_' + pred_ver].
        ↪clip(0,20)
        submission[['ID', 'item_cnt_month']].to_csv('ver%d_%s.csv' % (ver,
        ↪pred_ver), index = False)
```

0.283260703911113
0.3396436425815545

```
[51]: #RF, LGBM and XGBoost produced the lowest RMSE overall, so these will now be
        ↪extracted to build an ensemble model to predict
X_train_level2_v2 = X_train_level2[:, [0,1,3]]
X_test_level2_v2 = X_test_level2[:, [0,1,3]]
```

```
[54]: #Ensembling v2
pred_list = {}

#second level learning model via linear regression
print('Training Second level learning model via linear regression')

from sklearn.linear_model import (LinearRegression, SGDRegressor)

lr = LinearRegression()
lr.fit(X_train_level2_v2, y_train_level2)

#compute R-squared on train and test sets
# print('Train R-squared for %s is %f' % ('test_preds_lr_stacking',
    ↪sqrt(mean_squared_error(y_train_level2, lr.predict(X_train_level2))))))

test_preds_lr_stacking_v2 = lr.predict(X_test_level2_v2)
train_preds_lr_stacking_v2 = lr.predict(X_train_level2_v2)
print('Train RMSE for %s is %f' % ('train_preds_lr_stacking_v2',
    ↪sqrt(mean_squared_error(y_train_level2, train_preds_lr_stacking_v2))))

pred_list['test_preds_lr_stacking_v2'] = test_preds_lr_stacking_v2
if Validation:
    print('Test RMSE for %s is %f' % ('test_preds_lr_stacking_v2',
    ↪sqrt(mean_squared_error(y_test_level2, test_preds_lr_stacking_v2))))
```

Training Second level learning model via linear regression
Train RMSE for train_preds_lr_stacking_v2 is 0.747544

```
[55]: #second level learning model via SGDRegressor
print('Training Second level learning model via SGDRegressor')
sgdr = SGDRegressor(
    penalty = 'l2' ,
    random_state = SEED )

sgdr.fit(X_train_level2_v2, y_train_level2)

#compute R-squared on train and test sets
# print('Train R-squared for %s is %f' %('test_preds_lr_stacking',
    ↪sqrt(mean_squared_error(y_train_level2, lr.predict(X_train_level2))))
test_preds_sgdr_stacking_v2 = sgdr.predict(X_test_level2_v2)
train_preds_sgdr_stacking_v2 = sgdr.predict(X_train_level2_v2)
print('Train RMSE for %s is %f' %('train_preds_lr_stacking_v2',
    ↪sqrt(mean_squared_error(y_train_level2, train_preds_sgdr_stacking_v2))))

pred_list['test_preds_sgdr_stacking_v2'] = test_preds_sgdr_stacking_v2

if Validation:
    print('Test RMSE for %s is %f' %('test_preds_sgdr_stacking_v2',
    ↪sqrt(mean_squared_error(y_test_level2, test_preds_sgdr_stacking_v2))))

#print('%0.2f min: Finish training second level model'%((time.time() -
    ↪start_time)/60))
```

Training Second level learning model via SGDRegressor
Train RMSE for train_preds_lr_stacking_v2 is 0.755024

```
[56]: #submission
if not Validation:
    submission = pd.read_csv('sample_submission.csv')
    ver = 3

    for pred_ver in ['lr_stacking_v2', 'sgdr_stacking_v2']:
        print(pred_list['test_preds_' + pred_ver].clip(0,20).mean())
        submission['item_cnt_month'] = pred_list['test_preds_' + pred_ver].
    ↪clip(0,20)
        submission[['ID', 'item_cnt_month']].to_csv('ver%d_%s.csv' % (ver,
    ↪pred_ver), index = False)
```

0.281080262941557
0.32307098985203775

Despite having a higher RMSE for the training set, the version 2 SGDR ensemble performed the best on the test set when submitted to Kaggle.

Public and private LB scores for this project were: 0.950299 and 0.949603.

[]: